

HI!

CONCURRENCY: DATA STRUCTURES

Shivaram Venkataraman

CS 537, Spring 2019

ADMINISTRIVIA

Spring break!

AGENDA / LEARNING OUTCOMES

Concurrency: How to build concurrent data structures?

Summary of virtualization, concurrency

RECAP

CONCURRENCY OBJECTIVES

Mutual exclusion (e.g., A and B don't run at same time)

solved with locks

Ordering (e.g., B runs after A does something)

solved with condition variables and semaphores

Bugs → Deadlocks

ABSTRACTIONS

Concurrency

Objects, Lists, Hashtable

Semaphores

Locks, Condition variables

Atomic Primitives

OS
C library
Arch
Instructions

Spin locks

equally

* chg or test and
CAS

CONCURRENT DATA STRUCTURES

CONCURRENT DATA STRUCTURES

Counters 

Lists 

Hashtable 

Queues 

Start with a correct solution

Make it perform better!

Thread safety
→ Mutual exclusion
ordering

 *still*
correct

WHAT IS SCALABILITY

N times as much work on N cores as done on 1 core

Strong scaling

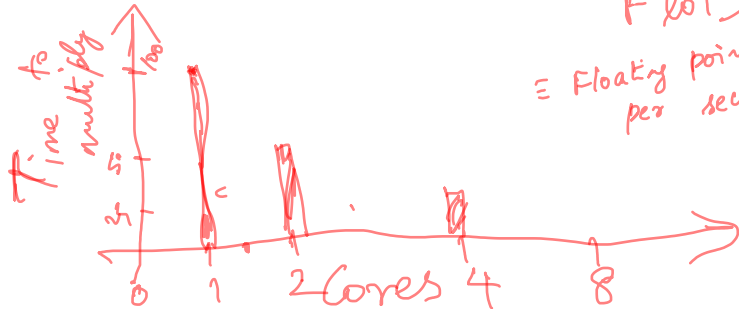
Fix input size, increase number of cores

e.g. $A \times B$
 $m \times n \quad n \times d$

$\equiv O(m \cdot n \cdot d)$
 FLOPS

\equiv Floating point ops per second

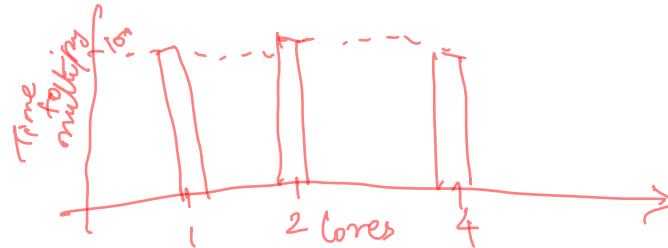
"Linear scaling"



Weak scaling

Increase input size with number of cores

1 core :	$100^A \times 100$	$100^B \times 100$	$10^A \times 10^B$
2 cores :	100×200	200×100	$2 \times 10^A \times 2 \times 10^B$
4 cores :	100×400	400×100	



COUNTERS

```
1 typedef struct __counter_t {  
2     int value;  
3 } counter_t;  
4  
5 void init(counter_t *c) {  
6     c->value = 0;  
7 }  
8 void increment(counter_t *c) {  
9     c->value++;  
10 }  
11 int get(counter_t *c) {  
12     return c->value;  
13 }
```

initialize

add

return

1

value

to zero

If two threads
increment

value shld increase
by 2

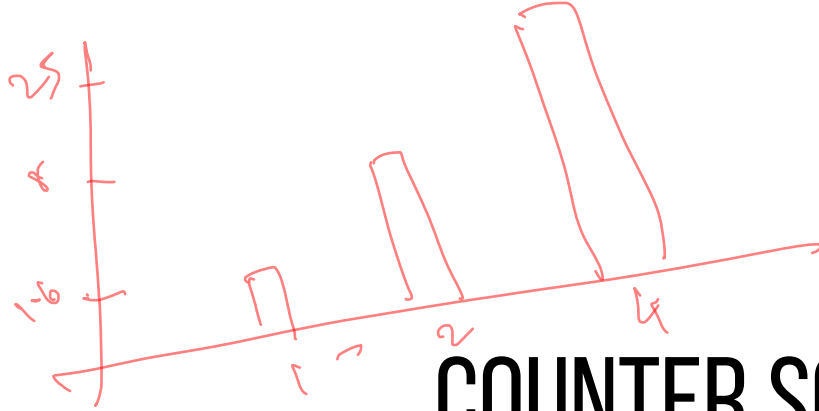
THREAD SAFE COUNTER

```
1 typedef struct __counter_t {  
2     int value;  
3     pthread_mutex_t lock;  
4 } counter_t;  
5  
...  
10  
11 void increment(counter_t *c) {  
12     pthread_mutex_lock(&c->lock);  
13     c->value++;  
14     pthread_mutex_unlock(&c->lock);  
15 }
```

CORRECT? → *How?*
T1 inc T2 inc T3 (get)

LINEARIZABILITY

```
int get(counter_t *c) {  
    pthread_mutex_lock(&c->lock);  
    c->value;  
    pthread_mutex_unlock(&c->lock);  
}
```



COUNTER SCALABILITY DEMO

UNDERLYING PROBLEM?

Ticket lock

```
void spin_lock(spinlock_t *lock)
{
    t = atomic_inc(lock->next_ticket);
    while (t != lock->current_ticket)
        ; /* Spin */
}
```

Turn

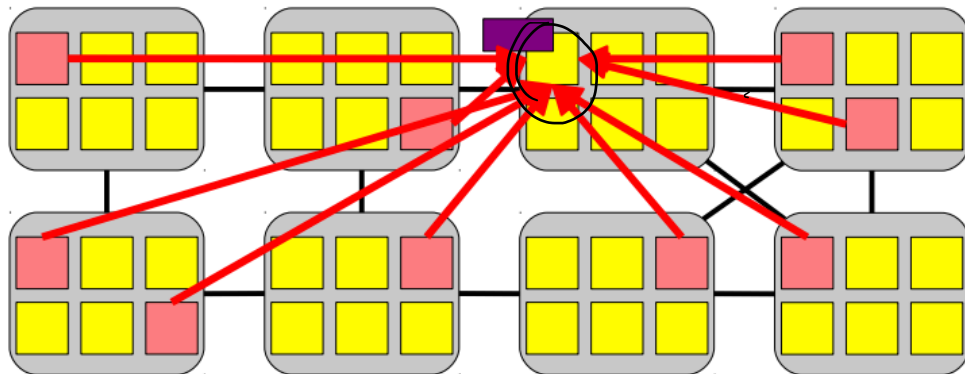
```
void spin_unlock(spinlock_t *lock)
{
    lock->current_ticket++;
}

struct spinlock_t {
    int current_ticket;
    int next_ticket;
}
```

up to 48 cores

An Analysis of Linux
Scalability
to Many Cores

Boyd-Wickizer et. al
OSDI 2010



APPROXIMATE COUNTERS

Maintain a counter per-core, global counter

Global counter lock

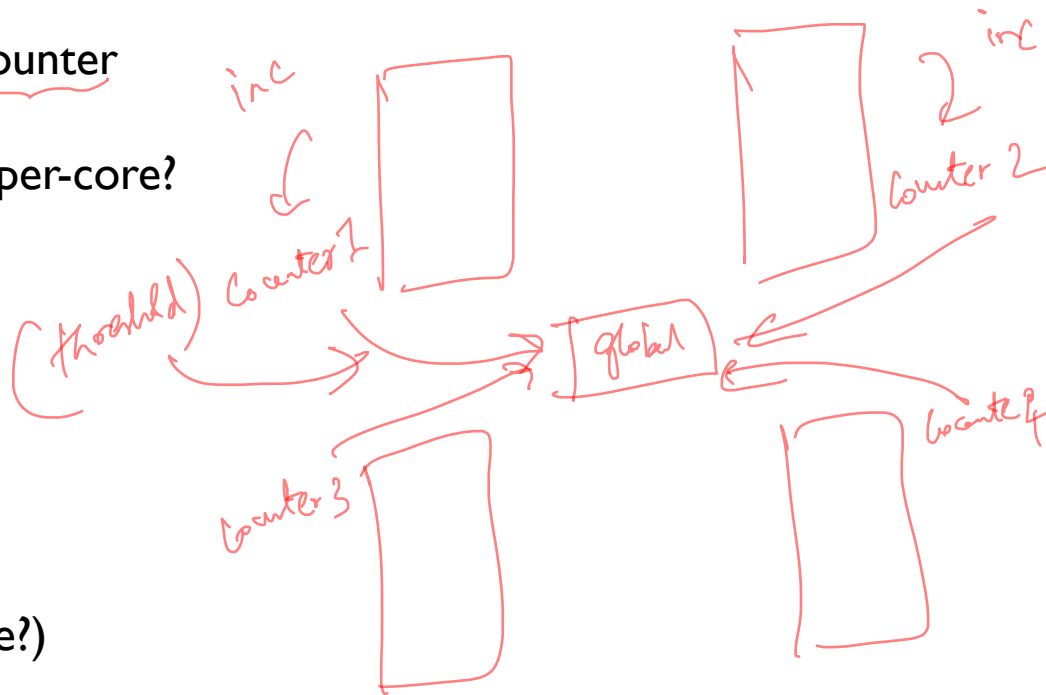
Per-core locks if more than 1 thread per-core?

Increment:

update local counters
at threshold update global

Read:

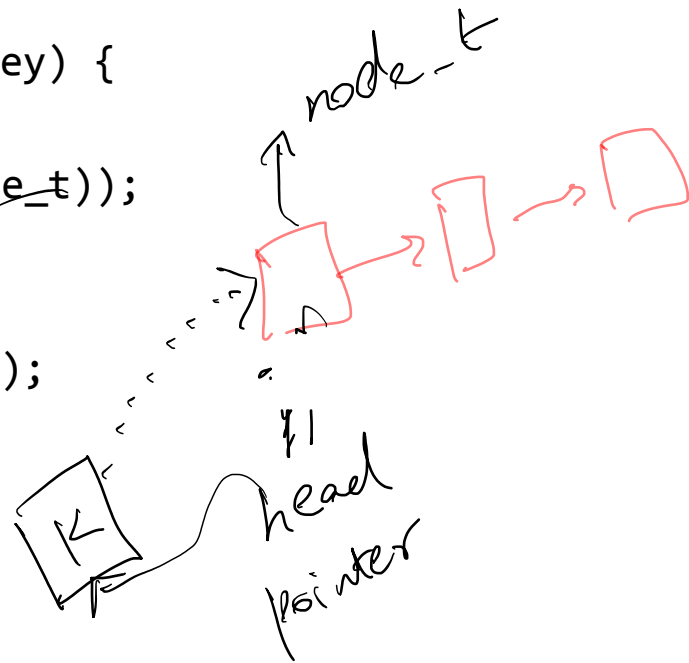
global counter (maybe inaccurate?)



DEMO

CONCURRENT LINKED LIST

```
18 void List_Insert(list_t *L, int key) {  
19     pthread_mutex_lock(&L->lock);  
20     node_t *new = malloc(sizeof(node_t));  
21     if (new == NULL) {  
22         perror("malloc");  
23         pthread_mutex_unlock(&L->lock);  
24         return; // fail  
25     }  
26     new->key = key;  
27     new->next = L->head;  
28     L->head = new;  
29     pthread_mutex_unlock(&L->lock);  
30     return; // success  
31 }
```



BETTER CONCURRENT LINKED LIST?

```
18 void List_Insert(list_t *L, int key) {
19     node_t *new = malloc(sizeof(node_t));
21     if (new == NULL) {
22         perror("malloc");
23         pthread_mutex_unlock(&L->lock);
24         return; // fail
25     }
26     new->key = key;
27     new->next = L->head;
28     L->head = new;
29     pthread_mutex_unlock(&L->lock);
30     return; // success
31 }
```

Thread safe

~~pthread_mutex_lock(&L->lock);~~

Keep your
critical
section
small

DEMO

HASH TABLE FROM LIST

```
1 #define BUCKETS (101)
2 typedef struct __hash_t {
3     list_t lists[BUCKETS];
4 } hash_t;
5
6 int Hash_Insert(hash_t *H, int key) {
7     int bucket = key % BUCKETS;
8     return List_Insert(&H->lists[bucket], key);
9 }
10
```

If you can avoid contention by using different locks, better scaling

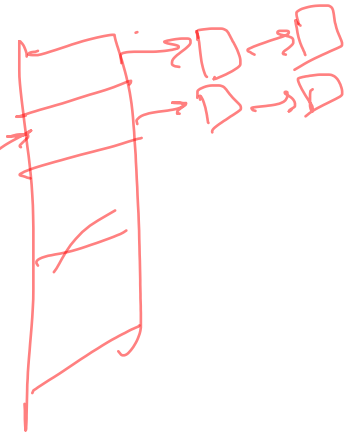
hashing?

hash(key) →

hashed key

lookup

value

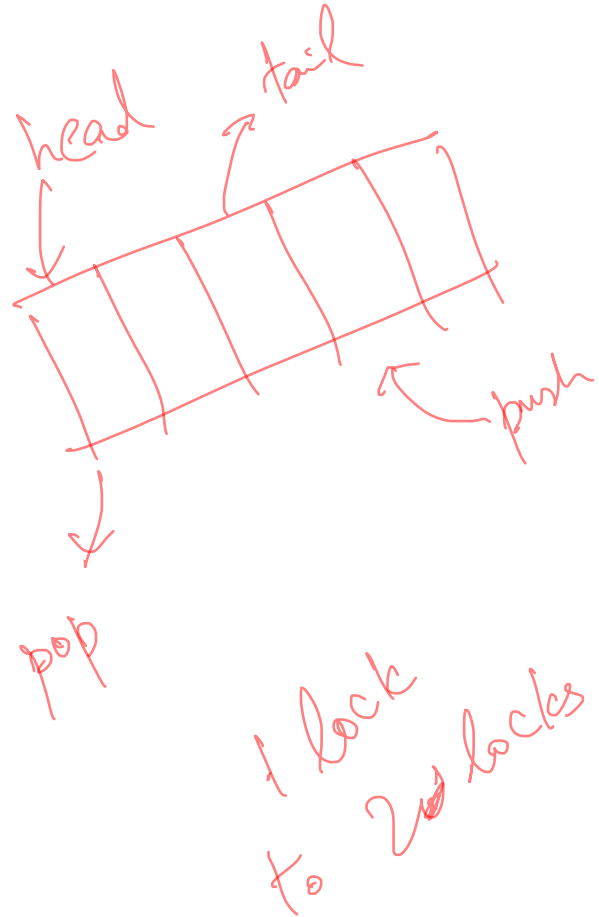


DEMO

```

21 void Queue_Enqueue(queue_t *q, int value) {
22     node_t *tmp = malloc(sizeof(node_t));
23     assert(tmp != NULL);
24     tmp->value = value;
25     tmp->next = NULL;
26
27     pthread_mutex_lock(&q->tailLock);
28     q->tail->next = tmp;
29     q->tail = tmp;
30     pthread_mutex_unlock(&q->tailLock);
31 }
32
33 int Queue_Dequeue(queue_t *q, int *value) {
34     pthread_mutex_lock(&q->headLock);
35     node_t *tmp = q->head;
36     node_t *newHead = tmp->next;
37     if (newHead == NULL) {
38         pthread_mutex_unlock(&q->headLock);
39         return -1; // queue was empty
40     }
41     *value = newHead->value;
42     q->head = newHead;
43     pthread_mutex_unlock(&q->headLock);
44     free(tmp);
45     return 0;
46 }

```



CONCURRENT DATA STRUCTURES

Simple approach: Add a lock to each method?!

Check for scalability – weak scaling, strong scaling

Avoid cross-thread, cross-core traffic

Per-core counter

Buckets in hashtable

Relaxed
Consistency

locks more
fine grained

Critical
sections
small

```
Java  
class Queue {  
    public synchronized get() {  
        < work >  
    }  
}
```

OPERATING SYSTEMS: THREE EASY PIECES

Three conceptual pieces



1. Virtualization

2. Concurrency



3. Persistence

VIRTUALIZATION

Make each application believe it has each **resource to itself**

CPU and Memory

Abstraction: Process API, Address spaces

Mechanism:

Limited direct execution, CPU scheduling

Address translation (segmentation, paging, TLB)

Policy: MLFQ, LRU etc.

Scheduler
memory allocation
stack growth

CONCURRENCY

Events occur simultaneously and may interact with one another

Need to

- Hide concurrency from independent processes

- Manage concurrency with interacting processes

Provide abstractions (locks, semaphores, condition variables etc.)

Correctness: mutual exclusion, ordering

Performance: scaling data structures, fairness

Common Bugs!

NEXT STEPS

Spring break!