

PERSISTENCE: I/O DEVICES

Shivaram Venkataraman

CS 537, Spring 2019

ADMINISTRIVIA

Project 4a: Out tonight, due on April 4th

Work in groups of up to two

Grades: Project 2b, 3, midterm by tomorrow!

AGENDA / LEARNING OUTCOMES

How does the OS interact with I/O devices?

What are the components of a hard disk drive?

RECAP

OPERATING SYSTEMS: THREE EASY PIECES

Three conceptual pieces

1. Virtualization

2. Concurrency

3. Persistence

VIRTUALIZATION

Make each application believe it has each resource to itself

CPU and Memory

Abstraction: Process API, Address spaces

Mechanism:

Limited direct execution, CPU scheduling

Address translation (segmentation, paging, TLB)

Policy: MLFQ, LRU etc.

CONCURRENCY

Events occur simultaneously and may interact with one another

Need to

Hide concurrency from independent processes

Manage concurrency with interacting processes

Provide abstractions (locks, semaphores, condition variables etc.)

Correctness: mutual exclusion, ordering

Performance: scaling data structures, fairness

Common Bugs!

T₁ after T₂
Producer Consumer queues

OPERATING SYSTEMS: THREE EASY PIECES

Three conceptual pieces

1. Virtualization

2. Concurrency

3. Persistence

MOTIVATION

What good is a computer without any I/O devices?

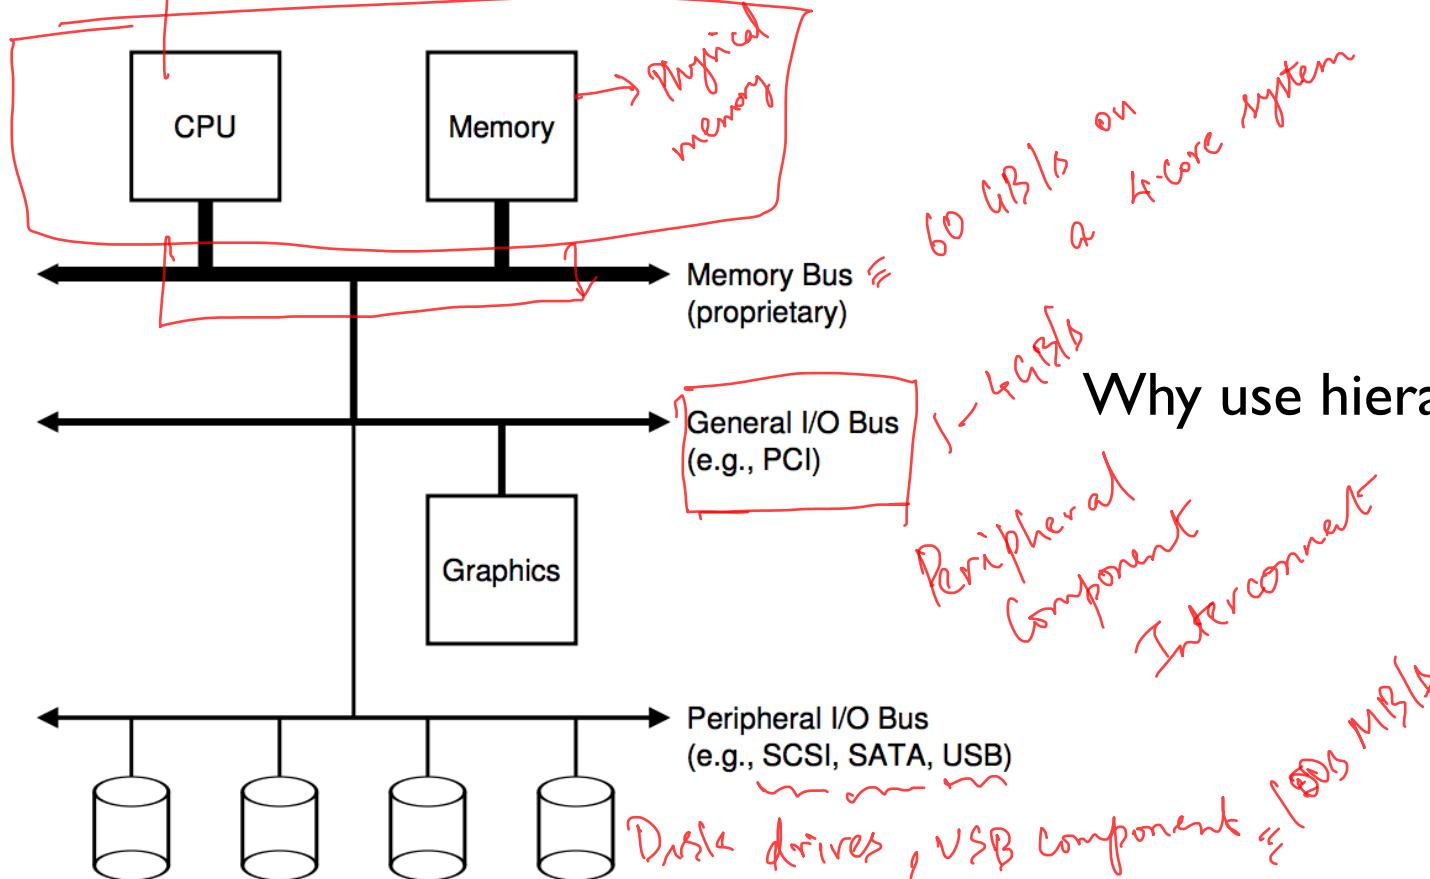
keyboard, display, disks

This per I/O devices


We want:

- **H/W** that will let us plug in different devices
- **OS** that can interact with different combinations

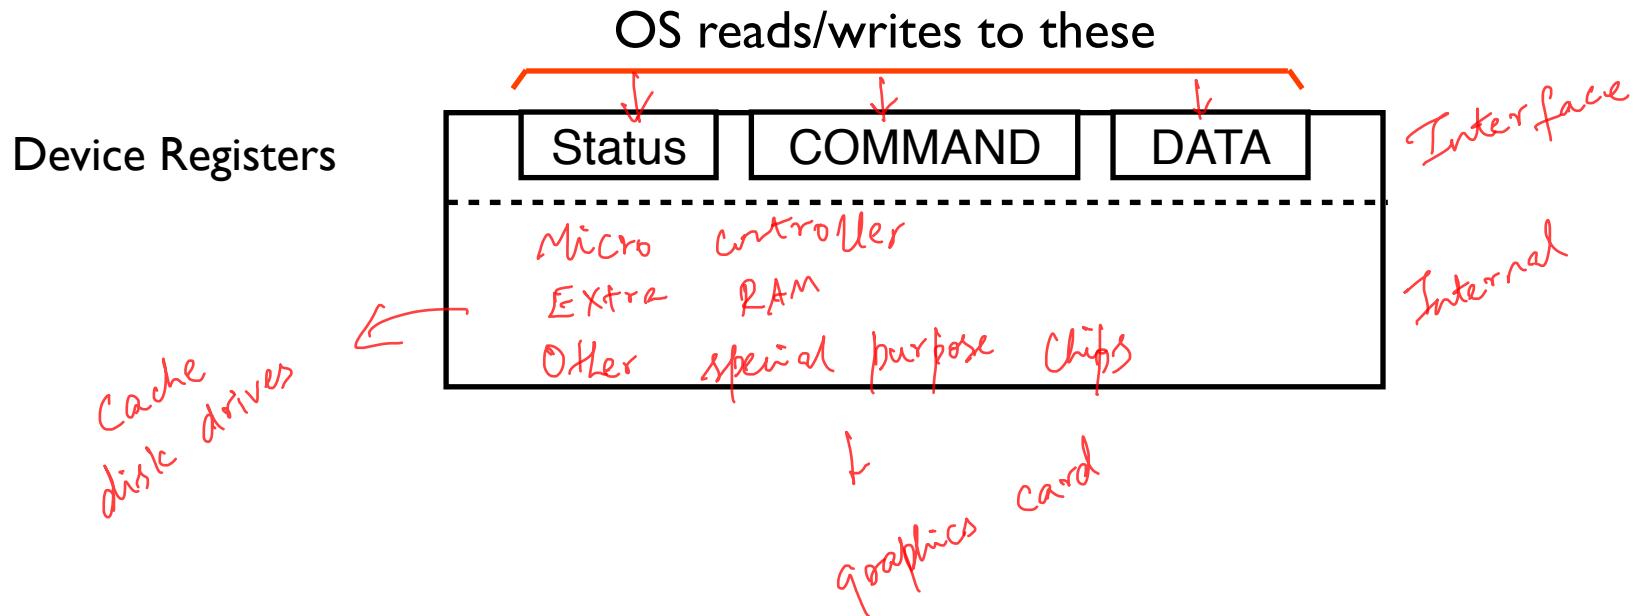
HARDWARE SUPPORT FOR I/O



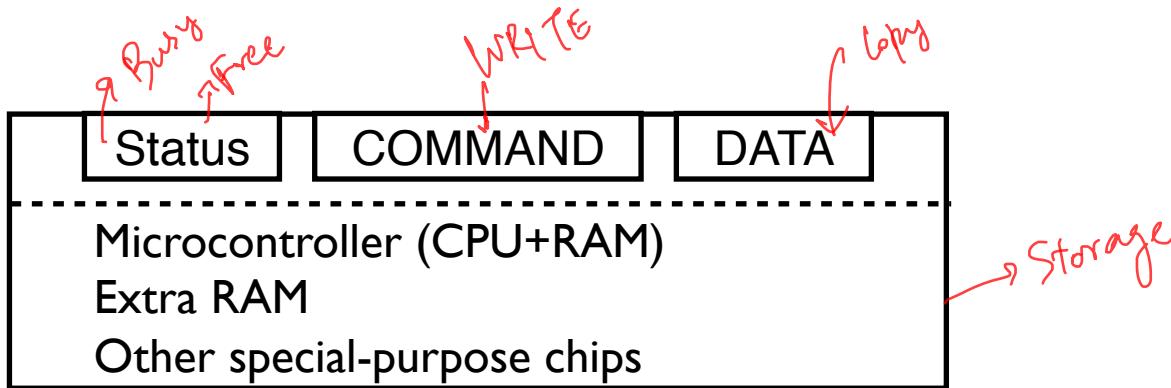
Why use hierarchical buses?



CANONICAL DEVICE



EXAMPLE WRITE PROTOCOL



- 1 while (STATUS == BUSY)
 ; // spin *Check if device is free*
- 2 Write data to DATA register
- 3 Write command to COMMAND register
- 4 while (STATUS == BUSY)
 ; // spin



```
lock(disk)
while (STATUS == BUSY)
;
```

// 1


Write data to DATA register // 2

Write command to COMMAND register // 3

while (STATUS == BUSY) // 4 //

;

unlock(disk)



Interrupts!

```
while (STATUS == BUSY) // 1
```

```
    wait for interrupt;
```

```
    Write data to DATA register
```

```
    Write command to COMMAND register
```

```
    while (STATUS == BUSY)
```

```
        wait for interrupt;
```

// 2


// 3

// 4

firelock
↳ no useful work

INTERRUPTS VS. POLLING

Are interrupts always better than polling?



Fast device: Better to spin than take interrupt overhead

- Device time unknown? Hybrid approach (spin then use interrupts)


Flood of interrupts arrive

↳ Context Switch $\approx 1\text{ms}$


- Can lead to livelock (always handling interrupts)
- Better to ignore interrupts while make some progress handling them

Other improvement

- Interrupt coalescing (batch together several interrupts)

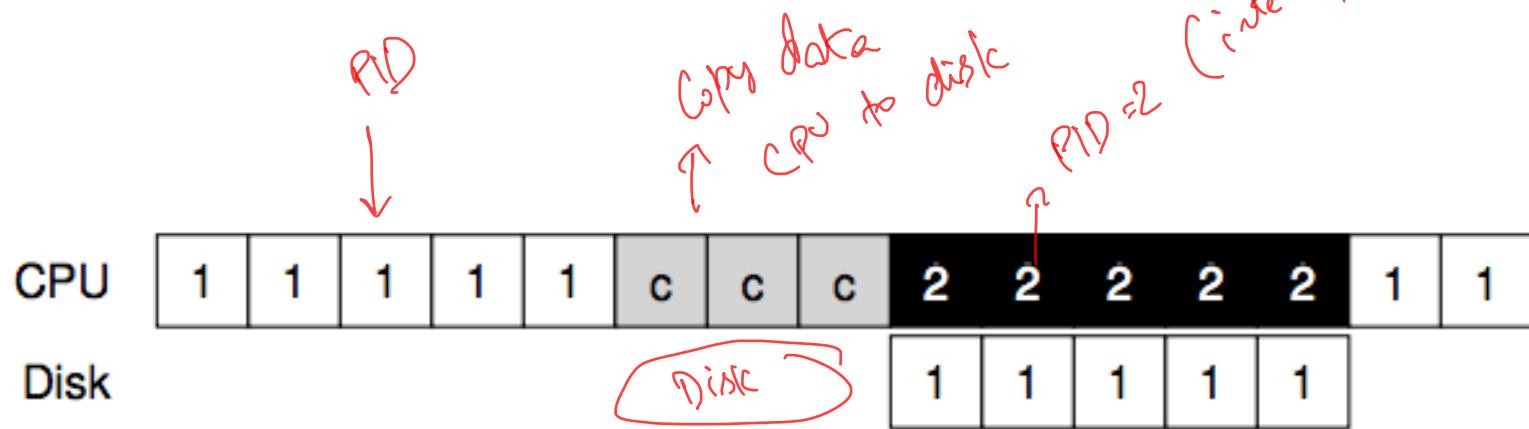


PROTOCOL VARIANTS



Status checks: polling vs. interrupts

DATA TRANSFER COSTS



Red handwritten annotation: "CPU is not used effectively"

PROGRAMMED I/O VS. DIRECT MEMORY ACCESS

DMA

PIO (Programmed I/O):

- CPU directly tells device what the data is


DMA (Direct Memory Access):

- CPU leaves data in memory
- Device reads data directly from memory

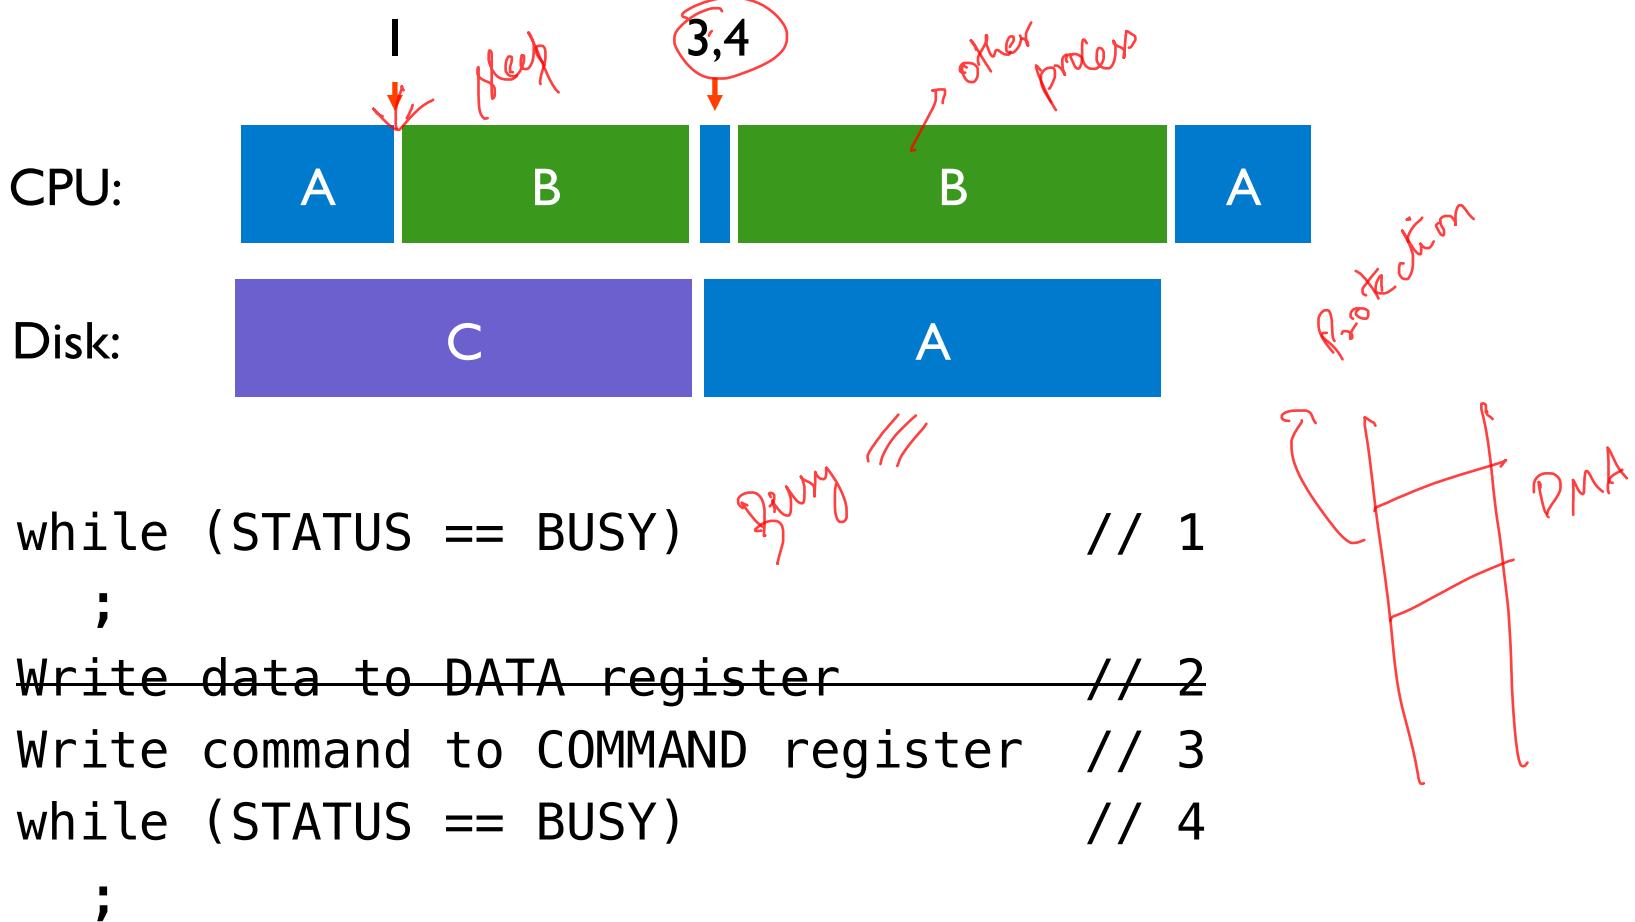


DMA

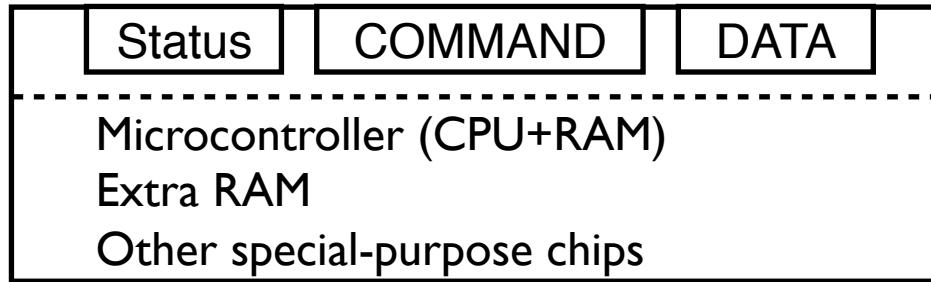
Disk



DMA engine

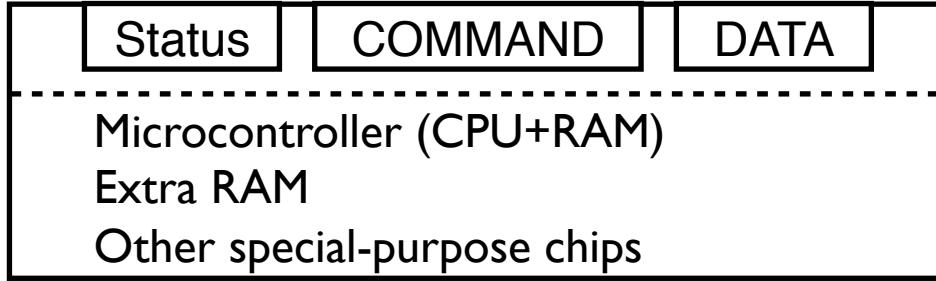


PROTOCOL VARIANTS



Status checks: polling vs. interrupts

PIO vs DMA



```
while (STATUS == BUSY)          // 1  
;  
Write data to DATA register    // 2  
Write command to COMMAND register // 3  
while (STATUS == BUSY)          // 4  
;
```

SPECIAL INSTRUCTIONS VS. MEM-MAPPED I/O

Special instructions


- each device has a port
 - in/out instructions (`x86`) communicate with device

xvii

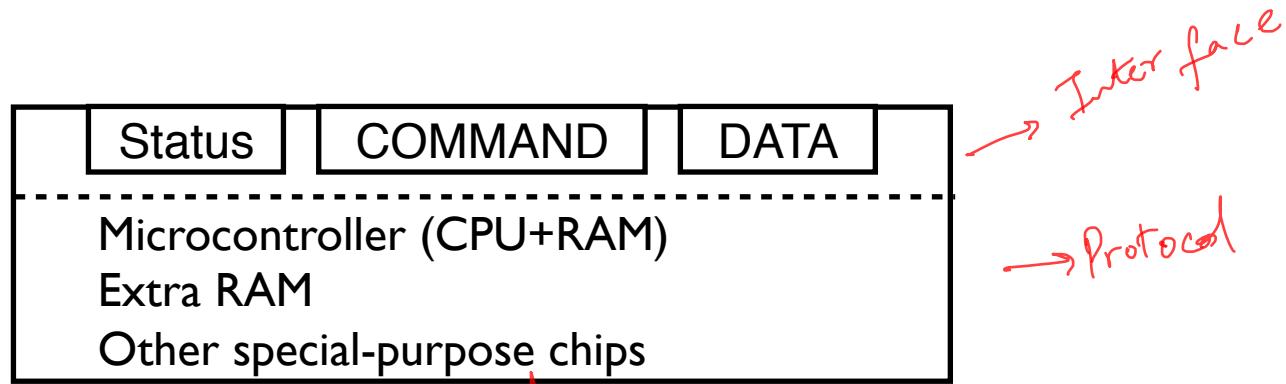
Memory-Mapped I/O

- H/W maps registers into address space
 - loads/stores sent to device

Doesn't matter much (both are used)



PROTOCOL VARIANTS



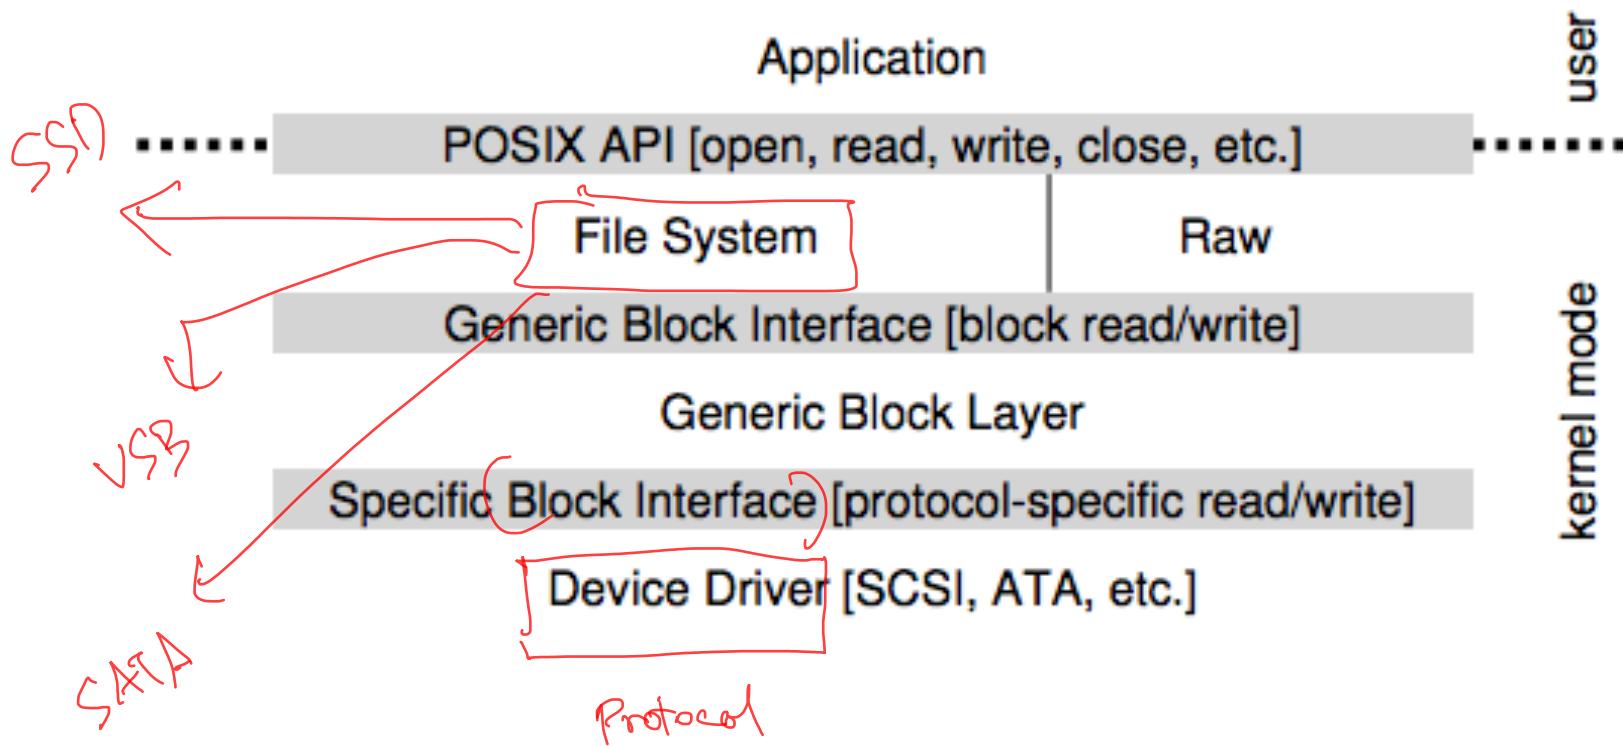
Status checks: polling vs. interrupts

PIO vs DMA

Special instructions vs. Memory mapped I/O

remove polling
DMA
In / out
memory mapped

DEVICE DRIVERS



VARIETY IS A CHALLENGE

Problem:

- many, many devices
- each has its own protocol

How can we avoid writing a slightly different OS for each H/W combination?

Write **device driver** for each device

Drivers are **70%** of Linux source code

Isolate
device driver

BUNNY 10



<https://tinyurl.com/cs537-sp19-bunny10>

BUNNY 10

SSD
volatile
device

new or
class
storage
devices

Intel Optane
drive

If you have a fast non-volatile memory based storage device, which approach would work better?

≈ few micro seconds

Polling

What part of a device protocol is improved by using DMA ?

Data register

HARD DISKS

HARD DISK INTERFACE


Disk has a sector-addressable address space
Appears as an array of sectors

Sectors are typically **512 bytes**


Main operations: reads + writes to sectors

Mechanical and slow (?)

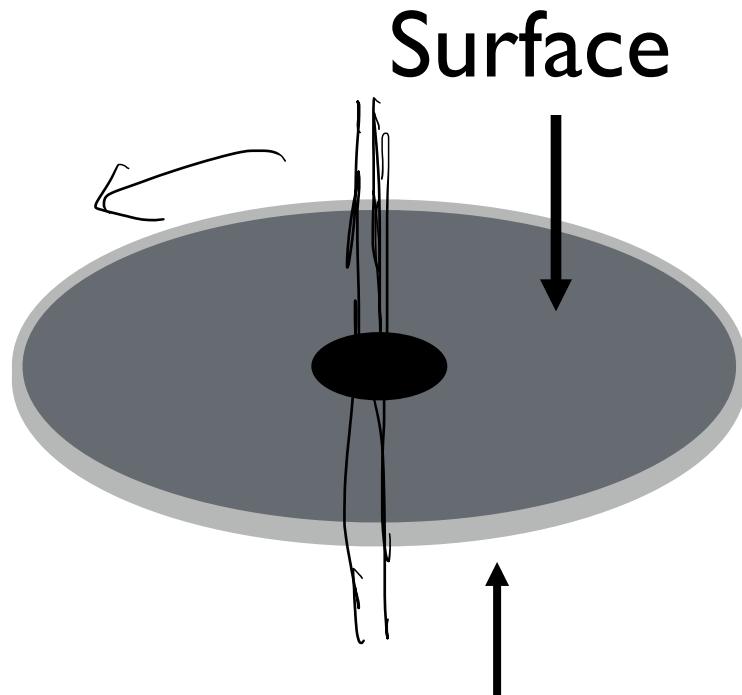
operations



Platter



Spindle



Surface

RPM?

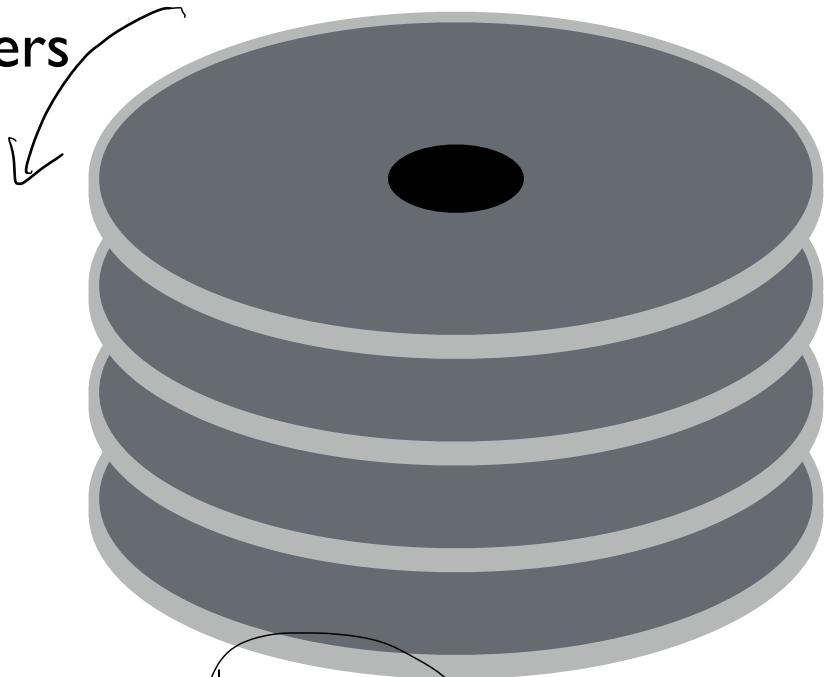
Motor connected to spindle **spins** platters

Rate of rotation: **RPM**

10000 RPM → single rotation is 6 ms

~~10000 rotations per min~~
~~10,000 rotations per sec~~

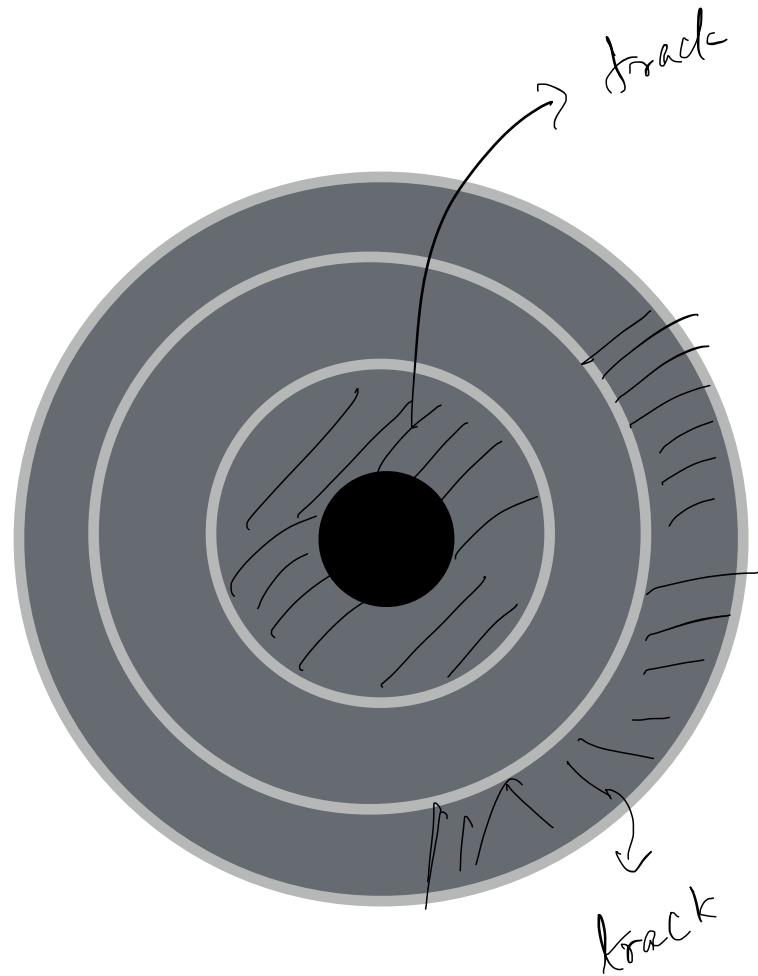
$1000 \text{ ms} \text{ in } 1 \text{ s}$
 $\Rightarrow \text{Latency} = \frac{60}{10,000} \times 1000$



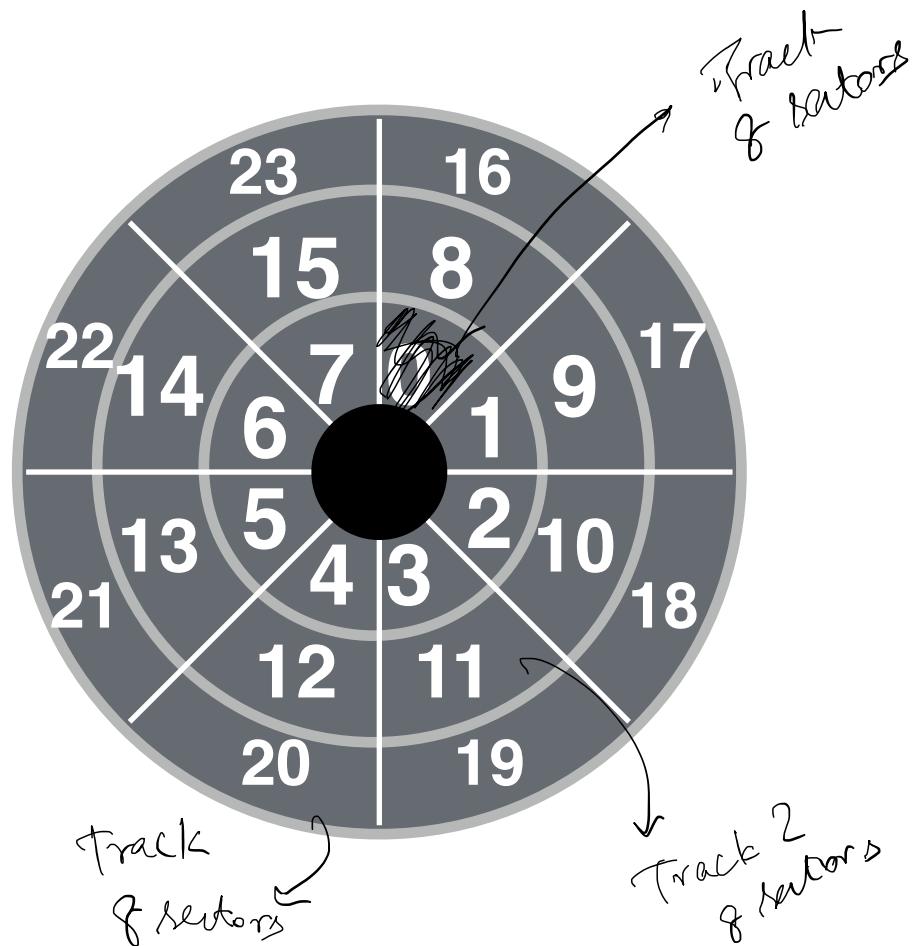
6 ms

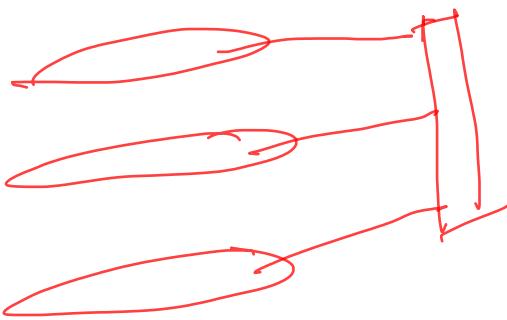
Surface is divided into rings: **tracks**

Stack of tracks(across platters): **cylinder**




Tracks are divided into numbered sectors






Heads on a moving **arm** can read from each surface.




READING DATA FROM DISK



Rotational delay

Average wait for a rotation half before sector reaches head

READING DATA FROM DISK



Seek Time

↓
Move the head right to track
the disk

TIME TO READ/WRITE

Three components:

Time = seek + rotation + transfer time



SEEK, ROTATE, TRANSFER



Seek cost: Function of cylinder distance

Not purely linear cost

Must accelerate, coast, decelerate, settle

Settling alone can take 0.5 - 2 ms

Entire seeks often takes 4 - 10 ms

Average seek = 1/3 of max seek

*Derivation
test hook*



Depends on rotations per minute (RPM)

7200 RPM is common, 15000 RPM is high end

Average rotation?

half of the rotation delay

Pretty fast: depends on RPM and sector density.

100+ MB/s is typical for maximum transfer rate

media , bus

BUNNY 11



<https://tinyurl.com/cs537-sp19-bunny11>

What is the time for 4KB random read?

$$= T_{\text{seek}} + T_{\text{rotation}} + T_{\text{transfer}}$$

Cheetah

$$T_{\text{seek}} = 4 \text{ ms}$$

$$T_{\text{rotation}} = 2 \text{ ms}$$

$$T_{\text{transfer}} = \frac{4 \text{ KB}}{125 \text{ MB/s}} = \frac{4 \text{ KB}}{125 \times 10^3 \text{ KB/s}}$$

$$\approx 30 \text{ ms} = \frac{4 \times 10^{-3}}{125} \frac{1}{s}$$

NEXT STEPS

Advanced disk features

Scheduling disk requests

Project 4a: Out tonight

Grades: Project 2b, 3, midterm by tomorrow!