

PERSISTENCE: DISK SCHEDULING

Shivaram Venkataraman

CS 537, Spring 2019

ADMINISTRIVIA

Grades: Project 2b, 3, midterm grades out!
See Piazza for regrade information

Project 4a is out! Due April 4th
More details in discussion section

Out of town Monday, Tue next week.
Guest lecture on Tuesday

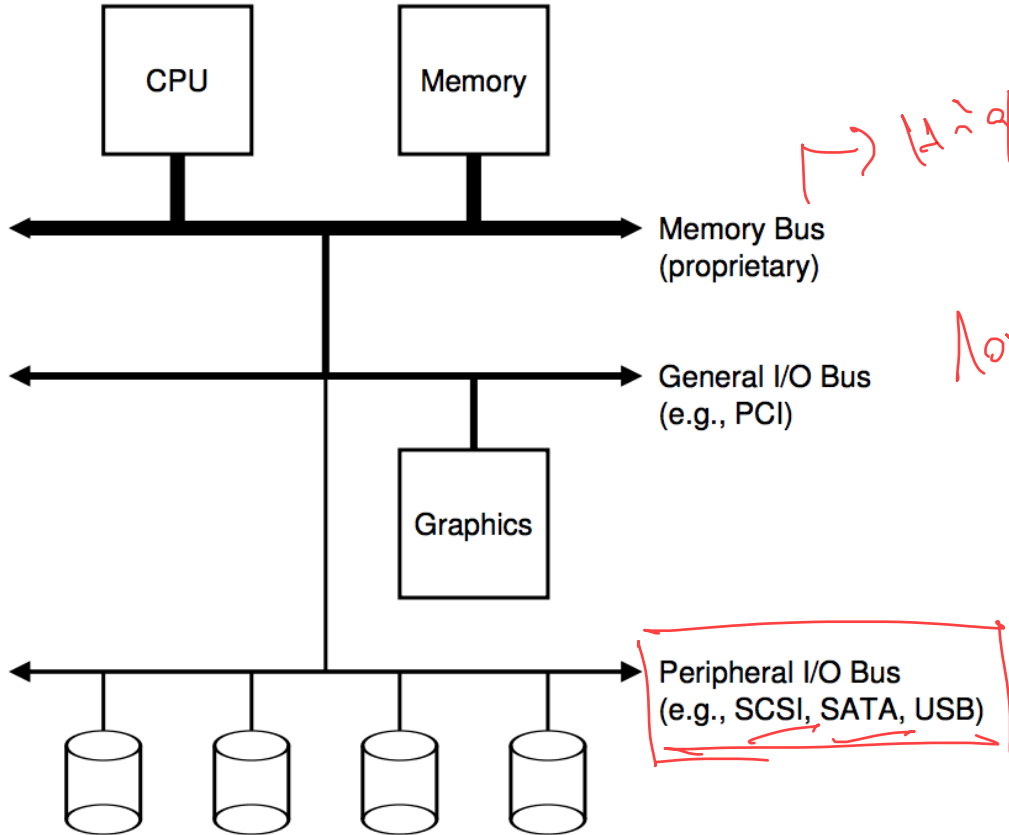
AGENDA / LEARNING OUTCOMES

How do you calculate sequential and random tput of a disk?

What algorithms are used to schedule I/O requests?

RECAP

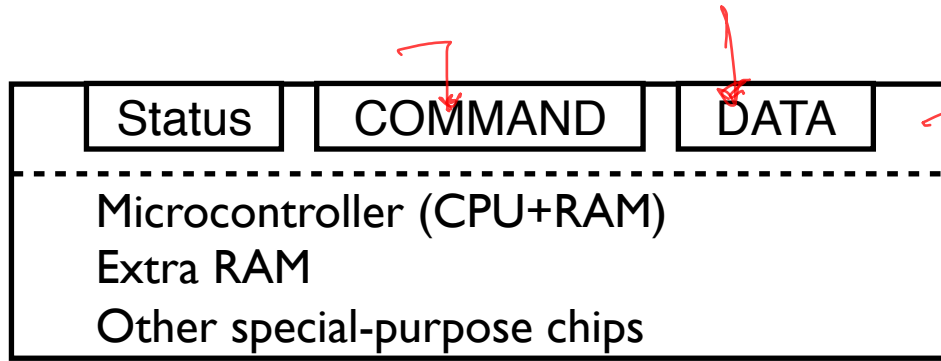
HARDWARE SUPPORT FOR I/O



High speed

Lower speed
greater distance
3.5 inch hard drive
SSD, keyboard

EXAMPLE WRITE PROTOCOL



waiting

```
while (STATUS == BUSY)  
    ; // spin
```

Copy Write data to DATA register

Write command to COMMAND register

waiting


```
while (STATUS == BUSY)  
    ; // spin
```

check if device is busy

wait for the operation

PROTOCOL VARIANTS

dev



Status	COMMAND	DATA
Microcontroller (CPU+RAM) Extra RAM Other special-purpose chips		

Status checks: polling vs. interrupts

CPU do other work

Trade-offs
very fast device \Rightarrow polling

PIO vs DMA

CPU stall avoided
"DMA"

Special instructions vs. Memory mapped I/O

IN/OUT

LOAD/STORES
special registers

programmable I/O

HARD DISK INTERFACE

Disk has a sector-addressable address space

Appears as an array of sectors


"Page" in memory

Sectors are typically **512 bytes**

Main operations: reads + writes to sectors

Mechanical and slow (?)

Atomic




RPM?


Motor connected to spindle **spins** platters

Rate of rotation: RPM


10000 RPM → single rotation is 6 ms



Heads on a moving **arm** can read from each surface.



READING DATA FROM DISK



1 Th disk
Num surface

sector size = 512 bytes

Seek Time \equiv

Rotational delay

Head right track
 \equiv Surface rotate right sector is below the Head

SEEK, ROTATE, TRANSFER

Seek cost: Function of cylinder distance
Not purely linear cost
Must accelerate, coast, decelerate, settle
Settling alone can take 0.5 - 2 ms

Entire seeks often takes 4 - 10 ms

Average seek = 1/3 of max seek

Depends on rotations per minute (RPM)
7200 RPM is common, 15000 RPM is high end

Average rotation?

Half the max rotation time

Pretty fast: depends on RPM and sector density.

100+ MB/s is typical for maximum transfer rate


BUNNY 11

What is the time for 4KB
random read?

	Cheetah 15K.5	Barracuda
Capacity	300 GB	1 TB
RPM	15,000	7,200
Average Seek	4 ms	9 ms
Max Transfer	125 MB/s	105 MB/s
Platters	4	4
Cache	16 MB	16/32 MB
Connects via	SCSI	SATA

High Perf
↑

Capacity
↑



How long does an average random 4-KB read take w/ Cheetah?

$$T_{I/O} = T_{seek} + T_{rotate} + T_{transfer}$$

	Cheetah 15K.5	Barracuda
Capacity	300 GB	1 TB
RPM	15,000	7,200
Average Seek	4 ms	9 ms
Max Transfer	125 MB/s	105 MB/s
Platters	4	4
Cache	16 MB	16/32 MB
Connects via	SCSI	SATA

$$T_{seek} = 4 \text{ ms}$$

$$T_{rotate} = \frac{1}{15000 \text{ RP}} \text{ ms} = \frac{1}{4} \text{ Rot per ms} \Rightarrow 4 \text{ ms / rot} \Rightarrow 2 \text{ ms on avg}$$

$$T_{transfer} = \frac{4 \text{ KB}}{125 \text{ MB}} \cdot 10^3 = \frac{4 \text{ KB}}{125 \times 10^3 \text{ KB}} \cdot 10^3 \text{ ms} = \frac{4}{125} \text{ ms} = 0.032 \text{ ms}$$

How long does an average random 4-KB read take w/ Barracuda?

$$T_{\text{seek}} = \boxed{9 \text{ ms}}$$

$$T_{\text{rot}} = \frac{7200 \text{ RPM}}{60 \times 1000} \text{ ms} = \frac{72}{600} \text{ s} = \frac{100}{672} \text{ ms/rot} = 8.33 \text{ ms}$$

Avg = $\boxed{4.16 \text{ ms}}$

$$T_{\text{transfer}} = \frac{4}{105} = \boxed{0.038 \text{ ms}}$$

$$= 13.2 \text{ ms}$$


	Cheetah 15K.5	Barracuda
Capacity	300 GB	1 TB
RPM	15,000	7,200
Average Seek	4 ms	9 ms
Max Transfer	125 MB/s	105 MB/s
Platters	4	4
Cache	16 MB	16/32 MB
Connects via	SCSI	SATA

WORKLOAD PERFORMANCE

WORKLOAD PERFORMANCE

So...


- seeks are slow
- rotations are slow
- transfers are fast



How does the kind of workload affect performance?

Sequential: access sectors in order

Random: access sectors arbitrarily



DISK SPEC

	Cheetah	Barracuda
Capacity	300 GB	1 TB
RPM	15,000	7,200
Avg Seek	4 ms	9 ms
Max Transfer	125 MB/s	105 MB/s
Platters	4	4
Cache	16 MB	32 MB

Sequential workload: what is throughput for each?

100MB
 Cheetah = $4\text{ms} + 2\text{ms} \approx 6\text{ms}$ $\Rightarrow 100\text{MB} / 6\text{ms} \approx 1,000$ effective $\approx 125\text{MB/s}$
 random 4KB = $6\text{ms} \Rightarrow 4\text{KB} / 6\text{ms} \lll 125\text{MB/s}$


OTHER IMPROVEMENTS

Track Skew


Zones

Cache

Imagine sequential reading,
how should sectors numbers be laid
out on disk?




When reading 16 after 15, the head won't settle quick enough, so we need to do a rotation.




Track Skew




Overlap & Seek
Trotate





outer tracks
more area



ZBR (Zoned bit recording): More sectors on outer tracks

OTHER IMPROVEMENTS


Track Skew

Zones

Cache

DRIVE CACHE

Drives may cache both reads and writes. (In addition to OS cache)




What advantage does caching in **drive** have for reads?

Store recently read sectors. | Fetch it from cache | Read ahead

What advantage does caching in **drive** have for writes?

CPU doesn't need to wait for write to finish. | Acknowledge before "persistent".



BUFFERING

Disks contain internal memory (2MB-16MB) used as cache

Read-ahead: “Track buffer”

- Read contents of entire track into memory during rotational delay

Write caching with volatile memory

- Immediate reporting: Claim written to disk when not
- Data could be lost on power failure

Tagged command queueing

- Have multiple outstanding requests to the disk
- Disk can reorder (schedule) requests for better performance


I/O SCHEDULERS

I/O SCHEDULERS

Given a stream of I/O requests, in what order should they be served?

Much different than CPU scheduling

Position of disk head relative to request position matters more than length of job





BUNNY12

<https://tinyurl.com/cs537-sp19-bunny12>

FCFS (FIRST-COME-FIRST-SERVE)

<https://tinyurl.com/cs537-sp19-bunny12>

FIFO

Assume seek+rotate = 10 ms for random request

How long (roughly) does the below workload take?

Requests are given in sector numbers

300001, 700001, 300002, 700002, 300003, 700003

60 ms

Why I/O scheduling
3x
⇒

300001, 300002, 300003, 700001, 700002, 700003

20 ms

SSTF (SHORTEST SEEK TIME FIRST)

Strategy always choose request that requires least seek time
(time for seeking and rotating)

next
best req
minimize
seek
30001

Greedy algorithm (just looks for best NEXT decision)

How to implement in OS?

↳ sector number "substitute"
nearest sector number

Disadvantages?

Starvation

30002, 30003

70001
starved

101, 102, 103, 104, ... 108, 201, ... 109

SCAN

Input 101, 201, 102, 301, 203

order 101, 201, 301, 203, 102

SCAN or Elevator Algorithm:




- Sweep back and forth, from one end of disk other, serving requests as pass that cylinder
- Sorts by cylinder number; ignores rotation delays

"Best effort" → "sorting"

C-SCAN (circular scan): Only sweep in one direction

Pros/Cons?

"Fair"



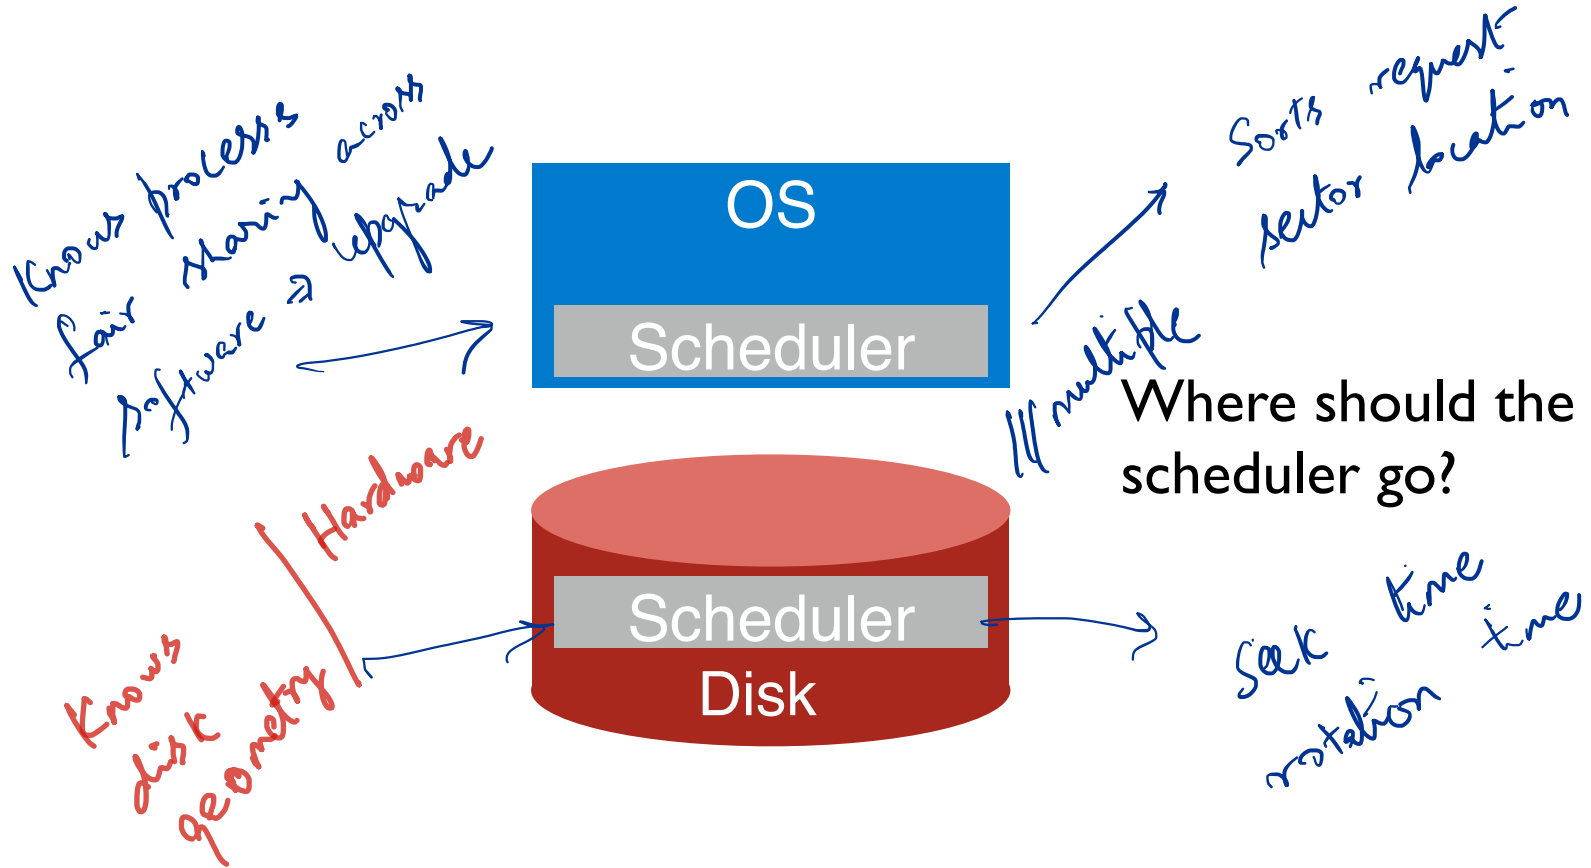
Rotates this way

Spindle

0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30 31

SATF (SHORTEST ACCESS TIME FIRST)



SCHEDULERS




WHAT HAPPENS?

Assume 2 processes each calling read() with C-SCAN

```
void reader(int fd) {  
    char buf[1024];  
    int rv;  
    while((rv = read(buf)) != 0) {  
        assert(rv);  
        // takes short time, e.g., 1ms  
        process(buf, rv);  
    }  
}
```



WORK CONSERVATION



Work conserving schedulers always try to do work if there's work to be done

Sometimes, it's better to wait instead if system anticipates another request will arrive

Possible improvements from I/O Merging



Work conserving
Always run
a request
if resource
is free

Not work
conserving
Wait for a
while
to "merge" or
get a better sequence

SUMMARY

Disks: Specific geometry with platters, spindle, tracks, sector

I/O Time: $\text{rotation_time} + \text{seek_time} + \text{transfer_time}$

Sequential throughput vs. random throughput

Advanced Techniques: Skewed layout, caching

Scheduling approaches: SSTF, SCAN, C-SCAN

Benefits of violating work conservation

DISK SIMULATOR HOMEWORK

Rotational speed is set to 1 degree per time. Complete revolution takes 360 time

Transfer begins and ends at the halfway point between sectors. E.g., to read sector 10, the transfer begins halfway between 9 and 10, ends halfway between 10 and 11.

There are 12 sectors per track, meaning that each sector takes up 30 degrees. To read a sector, it takes 30 time units (given our default speed of rotation).

Disk head is positioned on the outside track, halfway through sector 6.

Compute the seek, rotation, and transfer times for the following sets of requests:

1. -a 7
2. -a 7,30,8
3. -a 10,11,12,13

```
python disk.py -a <cmd> -G
```


Compare FIFO and SSTF for request stream 7,30,8

```
python disk.py -a 7,30,8 -p <SSTF|FIFO>
```

NEXT STEPS

Next class: How to achieve resilience against disk errors

Project 4a in Discussion today

Guest lecture on Tuesday