

Part A

1.

By simplification, the where condition becomes

(name = "John" or Name = "Mary") and (age = 20 or age > 50)

Give that name has only equality search, while age has both equality and range search, the five indexes could be as below:

1. Hash index on name, since hash support equality search
2. Clustered B+ tree index on age, since B+ tree support both equality and range search
3. Clustered B+ tree index on name, since B+ tree support equality search
4. Clustered B+ tree index on (name, age)
5. Clustered B+ tree index on (age, name)

2.

(a) Block Nested Loop Join:

$20000 + 42000 * \text{Ceil}(20000 / 9998) = 146000$ I/Os

(b) Sort-Merge Join

$B^2 = 10000^2 > 42000 = \text{Max}(42000, 20000) = \text{Max}(M_r, M_s)$, fits the optimal condition.

Since the data records are stored in the leaf nodes, we can save the cost of sorting. Therefore, we only need $20000 + 42000 = 62000$ I/Os

(c) Hash Join

$B^2 = 10000^2 > 20000 = \text{Min}(42000, 20000) = \text{Min}(M_r, M_s)$, fits the optimal condition.

Therefore, the cost is $3 * (20000 + 42000) = 186000$

=> Sort-Merge Join is the best one.

3.

City and business id are selected, which takes $20 + 8 = 28$ B.

Buffer pool size = $8 * 1024 * (10000 - 1) \geq 28 * f * \text{city\#}$ ($f = 1.4$)

$\text{city\#} \leq 2089586.94$

$\text{city\#} = 2089586$

4.

We can do a hash join users and reviews because $B^2 = 10000^2 > \text{Min}(M_r, M_s) = 75000$.

This costs $3 * (75000 + 500000) \text{ I/Os}$

We can then pipeline the result to selection and projection since we do not materialize the intermediate result

After selection and projection, then we can pipeline the result to the aggregation operator, which requires no I/O costs

Therefore, in total, the I/O cost is just the hash join cost which is
 $3 \times (75000 + 500000) = 1725000$ I/Os

5.

$\pi_{BName}((\pi_{BusinessID, BName} Businesses) \bowtie (\pi_{UserID}(\sigma_{ReviewCount \geq 100} Users)) \bowtie (\pi_{UserID, BusinessID}(\sigma_{Stars > 4} Reviews)))$

Part B

1.

Three Partition of R:

- 1) each of size N_R , number of partitions: $2B$
- 2) each of size $2N_R$, number of partitions: B
- 3) each of size $4N_R$, number of partitions: B

three sub-joining

Join S with partition 1) of R: $3(N_S + N_R) \times 2B = 6B \times N_S + 6B \times N_R$

Join S with partition 2) of R: $3(N_S + 2N_R) \times B = 3B \times N_S + 6B \times N_R$

Join S with partition 3) of R: $3(N_S + 4N_R) \times B = 3B \times N_S + 12B \times N_R$

Therefore, at the minimal, $12BN_S + 24BN_R$.

But we need to repartition sub-joining #3 because $2fN_R = 4B - 1$ and there are only $4B + 1$ buffer pages. $4fN_R$ must exceed the number of buffer pages.

I/O Cost for repartition – we only need to repartition B partitions of S and R respectively. To partition, we only need to consider one read and one write per page.

Therefore, it is $2 \times B \times (4N_R + N_S) = 8BN_R + 2BN_S$

In total, $12BN_S + 24BN_R + 8BN_R + 2BN_S = 14BN_S + 32BN_R$

2.

Using the idea of external merge sort, we sort elements in each run during pass 0.

In pass 1 of merging, we can compute the amount of each element. At final, we output the elements with the maximum amount. Therefore, the I/O cost will be the same as an external merge sort, which is

$$2N \left(\left\lceil \log_{B-1} \frac{N}{B} \right\rceil + 1 \right)$$

$$= 4 \times N = 4 \times 1000000 = 4000000 \text{ I/Os}$$