

ADVANCED SQL II

CS 564 - Fall 2018

ACKs: Dan Suciu, Jignesh Patel, AnHai Doan

WHAT IS THIS LECTURE ABOUT

- SQL: Aggregation
 - Aggregate operators
 - GROUP BY
 - HAVING
- SQL: Nulls
- SQL: Outer Joins

AGGREGATION

AGGREGATION

- **SUM, AVG, COUNT, MIN, MAX** can be applied to a column in a **SELECT** clause to produce that aggregation on the column
- **COUNT(*)** simply counts the number of tuples

```
SELECT AVG(Population)
FROM Country
WHERE Continent = 'Europe';
```

AGGREGATION: ELIMINATE DUPLICATES

We can use **COUNT(DISTINCT <attribute>)** to remove duplicate tuples before counting!

```
SELECT COUNT (DISTINCT Language)
FROM CountryLanguage ;
```

GROUP BY

- We may follow a **SELECT-FROM-WHERE** expression by **GROUP BY** and a list of attributes
- The relation is then grouped according to the values of those attributes, and any aggregation is applied only **within each group**

```
SELECT Continent, COUNT(*)  
FROM Country  
GROUP BY Continent;
```

GROUP BY: EXAMPLE

```
SELECT A, SUM(B * C)
FROM R
GROUP BY A;
```

R

A	B	C
a	2	0
a	5	1
b	7	1
b	6	0
c	4	1

grouping

A	B	C
a	2	0
	5	1
b	7	1
	6	0
c	4	1

select
clause

$$5 = 2*0 + 5*1$$

A	SUM(B*C)
a	5
b	7
c	4

RESTRICTIONS

If any aggregation is used, then each element of the **SELECT** list must be either:

- aggregated, or
- an attribute on the **GROUP BY** list

This query is **wrong!!**

```
SELECT Continent, COUNT(Code)
FROM Country
GROUP BY Code;
```


GROUP BY + HAVING

- The **HAVING** clause **always** follows a **GROUP BY** clause in a SQL query
 - it applies to each group, and groups not satisfying the condition are removed
 - it can refer only to attributes of relations in the **FROM** clause, as long as the attribute makes sense within a group

The HAVING clause applies **only** on aggregates!

HAVING: EXAMPLE

```
SELECT Language, COUNT(CountryCode) AS N
FROM CountryLanguage
WHERE Percentage >= 50
GROUP BY Language
HAVING N > 2
ORDER BY N DESC ;
```

PUTTING IT ALL TOGETHER

```
SELECT [DISTINCT] S
FROM R, S, T ,...
WHERE C1
GROUP BY attributes
HAVING C2
ORDER BY attribute ASC/DESC
LIMIT N ;
```

CONCEPTUAL EVALUATION

1. Compute the **FROM-WHERE** part, obtain a table with all attributes in R,S,T,...
2. Group the attributes in the **GROUP BY**
3. Compute the aggregates and keep only groups satisfying condition **C2** in the **HAVING** clause
4. Compute aggregates in S
5. Order by the attributes specified in **ORDER BY**
6. Limit the output if necessary

NULL VALUES

NULL VALUES

- tuples in SQL relations can have **NULL** as a value for one or more attributes
- The meaning depends on context:
 - **Missing value**: *e.g.* we know that Greece has some population, but we don't know what it is
 - **Inapplicable**: *e.g.* the value of attribute *spouse* for an unmarried person

NULL PROPAGATION

- When we do arithmetic operations using **NULL**, the result is again a **NULL**
 - $(10 * x) + 5$ returns **NULL** if $x = \text{NULL}$
 - **NULL**/0 also returns **NULL**!
- String concatenation also results in **NULL** when one of the operands is **NULL**
 - 'Wisconsin' || **NULL** || '-Madison' returns **NULL**

COMPARISONS WITH NULL

- The logic of conditions in SQL is **3-valued logic**:
 - **TRUE** = 1
 - **FALSE** = 0
 - **UNKNOWN** = 0.5
- When any value is compared with a **NULL**, the result is **UNKNOWN**
 - *e.g.* $x > 5$ is **UNKNOWN** if $x = \text{NULL}$
- A query produces a tuple in the answer **only if** its truth value in the **WHERE** clause is **TRUE** (1)

3-VALUED LOGIC

The truth value of a **WHERE** clause is computed using the following rules:

- **C1 AND C2** ----> $\min\{\text{value}(C1), \text{value}(C2)\}$
- **C1 OR C2** ----> $\max\{\text{value}(C1), \text{value}(C2)\}$
- **NOT C** ----> $1 - \text{value}(C)$

3-VALUED LOGIC: EXAMPLE

SELECT *

FROM R

WHERE $(R.A > 0)$ AND $((R.B < 5)$ OR $(\text{NOT } R.C = 3))$;

$\underbrace{\hspace{1.5cm}}$

1

$\underbrace{\hspace{1.5cm}}$

0.5

$\underbrace{\hspace{1.5cm}}$

0.5

$\underbrace{\hspace{2.5cm}}$

0.5 (1-0.5)

$\underbrace{\hspace{4.5cm}}$

0.5 (max{0.5, 0.5})

$\underbrace{\hspace{6.5cm}}$

0.5 (min{0.5, 1})

tuple (1, NULL, NULL)

the expression is **UNKNOWN!**

COMPLICATIONS

What will happen in the following query?

```
SELECT COUNT(*)  
FROM Country  
WHERE IndepYear > 1990 OR IndepYear <= 1990 ;
```

It will not count the rows with NULL!

TESTING FOR NULL

We can test for **NULL** explicitly:

- **x IS NULL**
- **x IS NOT NULL**

```
SELECT COUNT(*)  
FROM Country  
WHERE IndepYear > 1990 OR IndepYear <= 1990  
OR IndepYear IS NULL;
```

OUTER JOINS

INNER JOINS

The joins we have seen so far are **inner joins**

```
SELECT C.Name AS Country, MAX(T.Population) AS N
FROM Country C, City T
WHERE C.Code = T.CountryCode
GROUP BY C.Name;
```

Alternative syntax:

```
SELECT C.Name AS Country, MAX(T.Population) AS N
FROM Country C
INNER JOIN City T ON C.Code = T.CountryCode
GROUP BY C.Name;
```

We can simply also write **JOIN**

LEFT OUTER JOINS

A **left outer join** includes tuples from the left relation even if there's no match on the right! It fills the remaining attributes with NULL

```
SELECT C.Name AS Country, MAX(T.Population)
FROM Country C
LEFT OUTER JOIN City T
    ON C.Code = T.CountryCode
GROUP BY C.Name ;
```

LEFT OUTER JOIN: EXAMPLE

**SELECT A, C
FROM R LEFT OUTER JOIN S
ON R.B = S.B**

R

A	B
a	2
a	5
b	5
c	6

S

B	C
2	100
3	200
5	300
7	400



A	C
a	100
a	300
b	300
c	NULL

OTHER OUTER JOINS

- **Left outer join:**
 - include the left tuple even if there is no match
- **Right outer join:**
 - include the right tuple even if there is no match
- **Full outer join:**
 - include the both left and right tuples even if there is no match