

# HASH INDEXES

---

*CS 564- Fall 2018*

---

*ACKs: Dan Suciu, Jignesh Patel, AnHai Doan*

---

# WHAT IS THIS LECTURE ABOUT?

---

## Hash indexes

- Static Hashing
  - what is the I/O cost?
  - problems with static hashing
- Extendible Hashing
  - insertion
  - deletion

---

# HOW TO EVALUATE AN INDEX?

---

- What **access types** does it support?
  - *e.g.* equality search, range search, etc.
- Time to **access** a record
- Time to **insert** a record
- Time to **delete** a record
- How much **space** does it use?

# HASH INDEXES

---

- efficient for equality search
- not appropriate for range search
- Types of hash indexes:
  - static hashing
  - extendible (dynamic) hashing

# STATIC HASHING

---


- A hash index is a collection of *buckets*
  - bucket = primary page + overflow pages
  - each bucket contains one or more data entries
- To find the bucket for each record, we use a hash function  $h$  applied on the search key  $k$ 
  - $N$  = number of buckets
  - $h(k) \bmod N$  = bucket in which the data entry belongs
- Records with different search key may belong in the same bucket

# STATIC HASHING: EXAMPLE

**Person**(name, zipcode, phone)

- *search key*: zipcode
- *hash function  $h$* : last 2 digits

- 4 buckets
- each bucket has 2 data entries (full record)



# OPERATIONS ON HASH INDEXES

## Equality search (*search-key = value*)

- apply the hash function on the search key to locate the appropriate bucket
- search through the primary page (plus overflow pages) to find the record(s)

$$\text{I/O cost} = 1 + \text{\#overflow pages}$$

---

# OPERATIONS ON HASH INDEXES

---

- **Deletion**
  - find the appropriate bucket, delete the record
- **Insertion**
  - find the appropriate bucket, insert the record
  - if there is no space, create a new overflow page



# HASH FUNCTIONS

---

- An *ideal* hash function must be **uniform**: each bucket is assigned the same number of key values
- A *bad* hash function maps all search key values to the same bucket
- Examples of good hash functions:
  - $h(k) = a * k + b$ , where  $a$  and  $b$  are constants
  - a random function

---

# BUCKET OVERFLOW

---

- Bucket *overflow* can occur because of
  - insufficient number of buckets
  - *skew* in distribution of records
    - many records have the same search-key value
    - the hash function results in a non-uniform distribution of key values
- Bucket overflow is handled using *overflow buckets*

---

# PROBLEMS OF STATIC HASHING

---

- In static hashing, there is a **fixed** number of buckets in the index
- Issues with this:
  - if the database grows, the number of buckets will be too small: long overflow chains degrade performance
  - if the database shrinks, space is wasted
  - reorganizing the index is expensive and can block query execution

---

# EXTENDIBLE HASHING

---

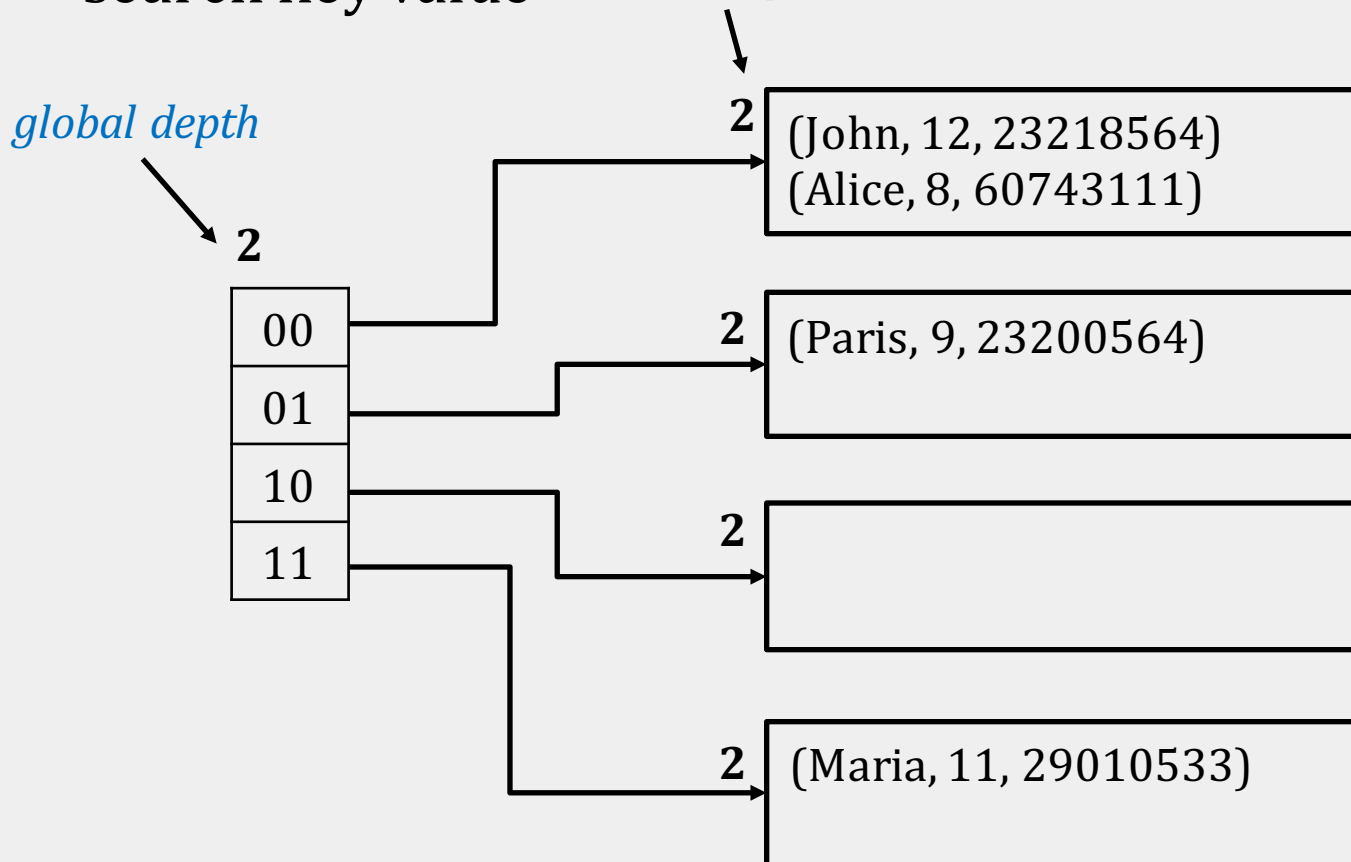
# EXTENDIBLE HASHING

---

- **Extendible hashing** is a type of *dynamic* hashing
- It keeps a directory of pointers to buckets
- On overflow, it reorganizes the index by **doubling the directory** (and not the number of buckets)

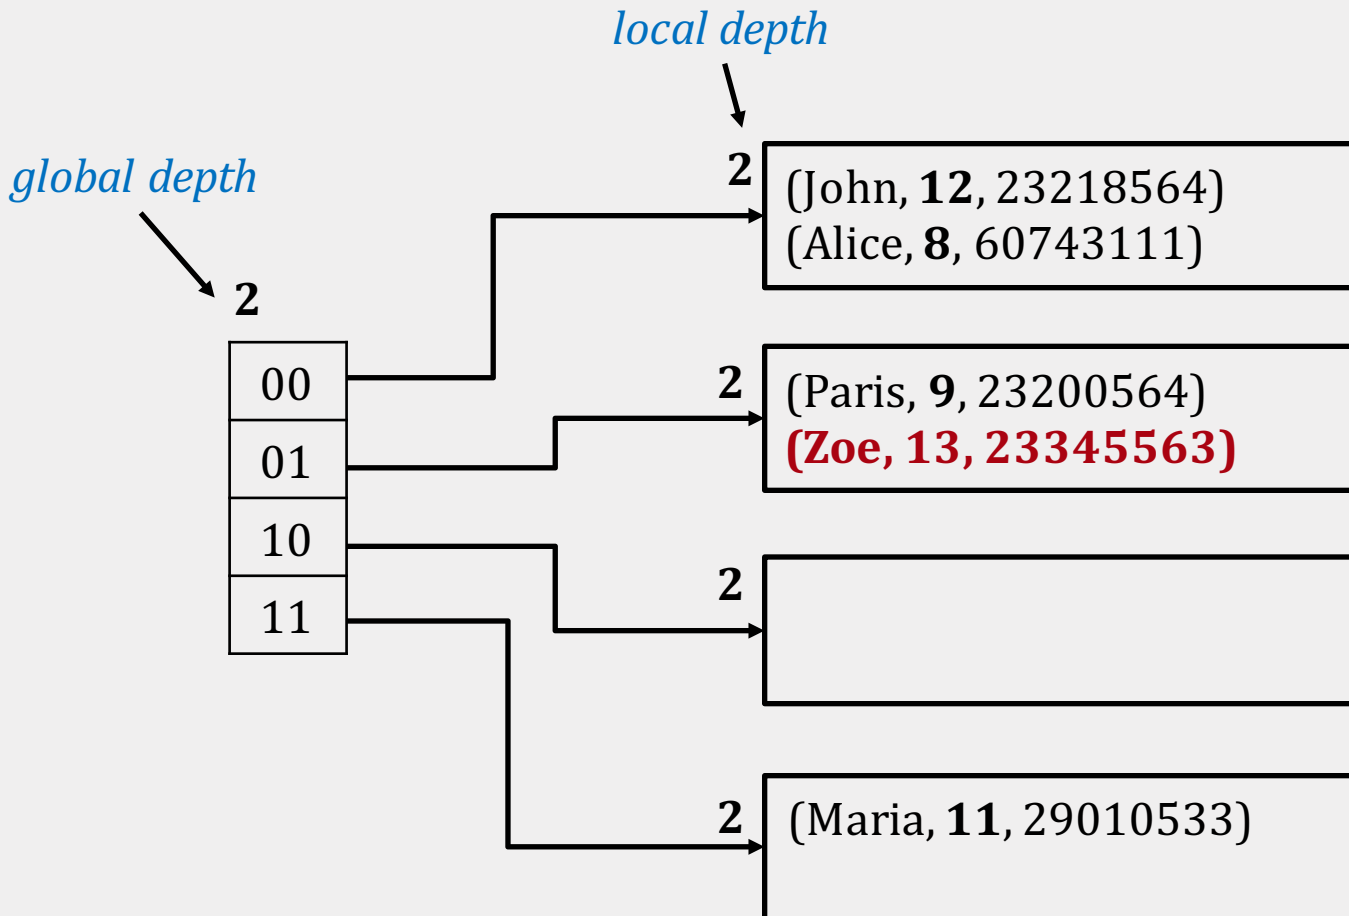
# EXTENDIBLE HASHING

To search, use the last **2** digits of the **binary** form of the search key value *local depth*



# EXTENDIBLE HASHING: INSERT

If there is space in the bucket, simply add the record

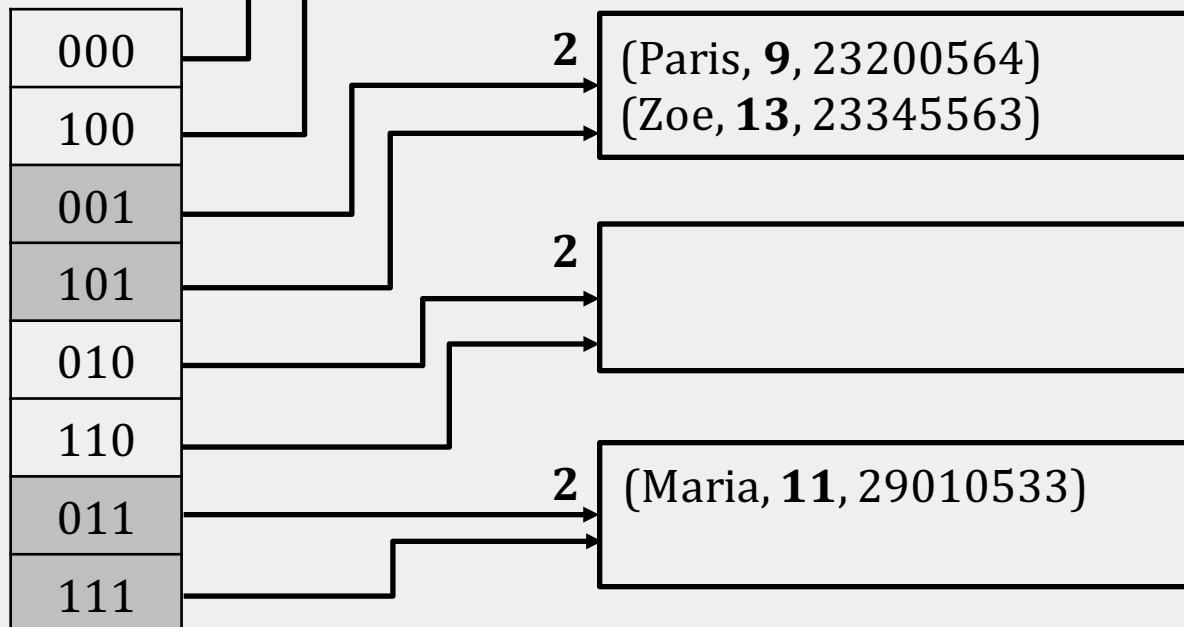


# EXTENDIBLE HASHING: INSERT

If the bucket is full, split the bucket and redistribute the entries

*global depth increases by 1*

3



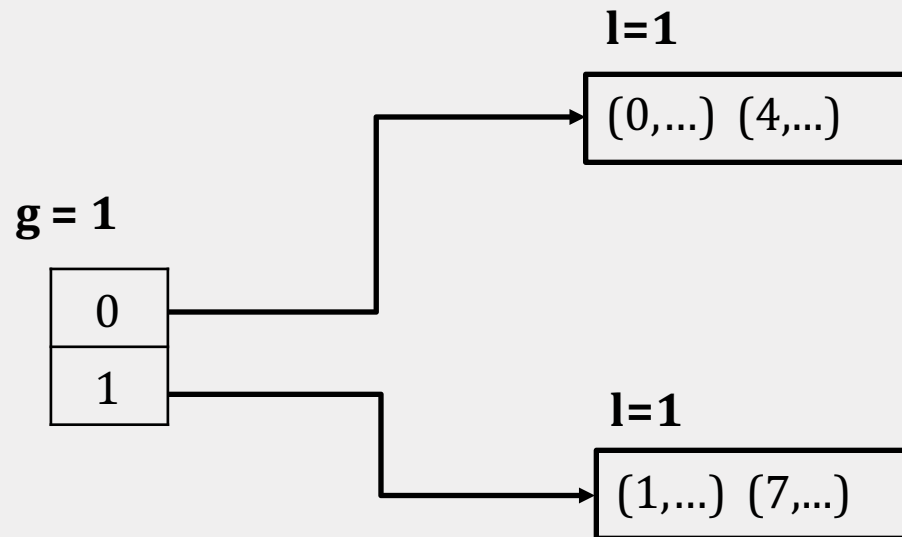
*local depth increases for the split bucket!*

*local depth remains the same for the other buckets*



# EXAMPLE

each page can hold at most two records

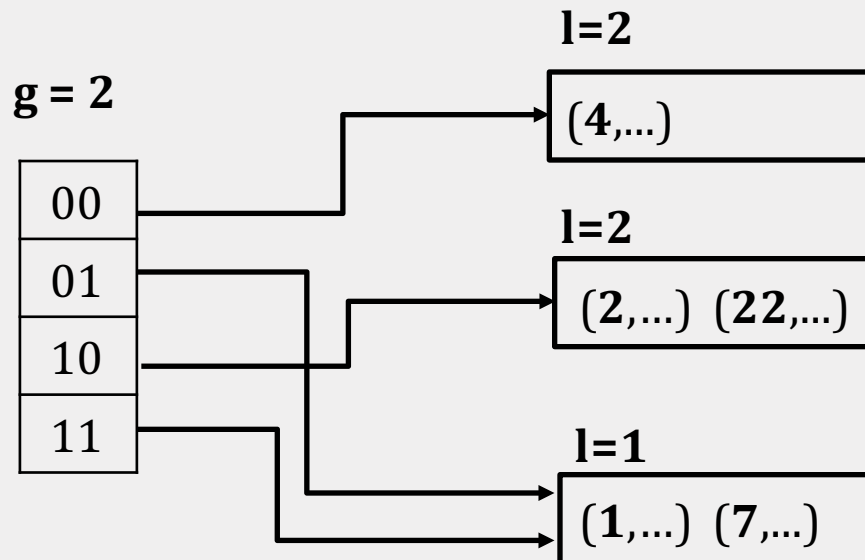


We always have: global depth  $\geq$  local depth

# EXAMPLE

- The catalog doubles in size
- Global depth becomes 2

insert: (22,...)



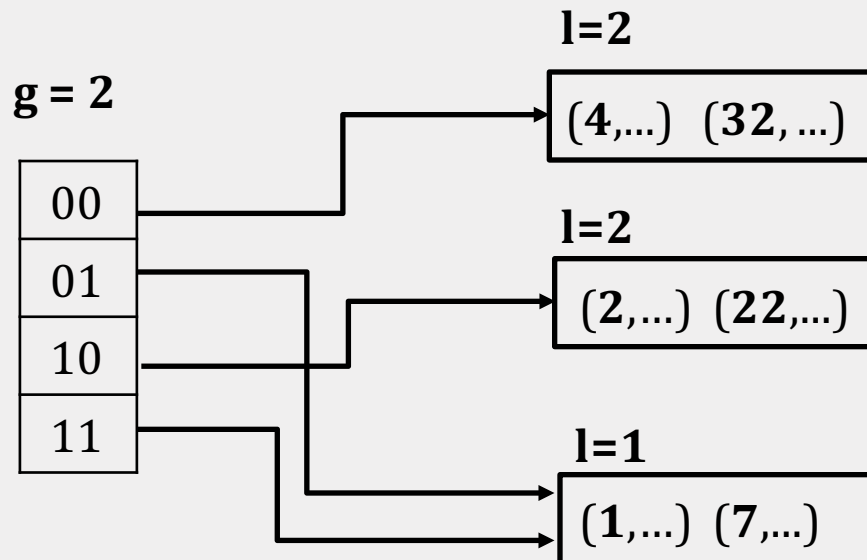
The bucket is split into two buckets with local depth 2

This bucket remains the same

# EXAMPLE

There is space in the bucket  
so nothing changes!

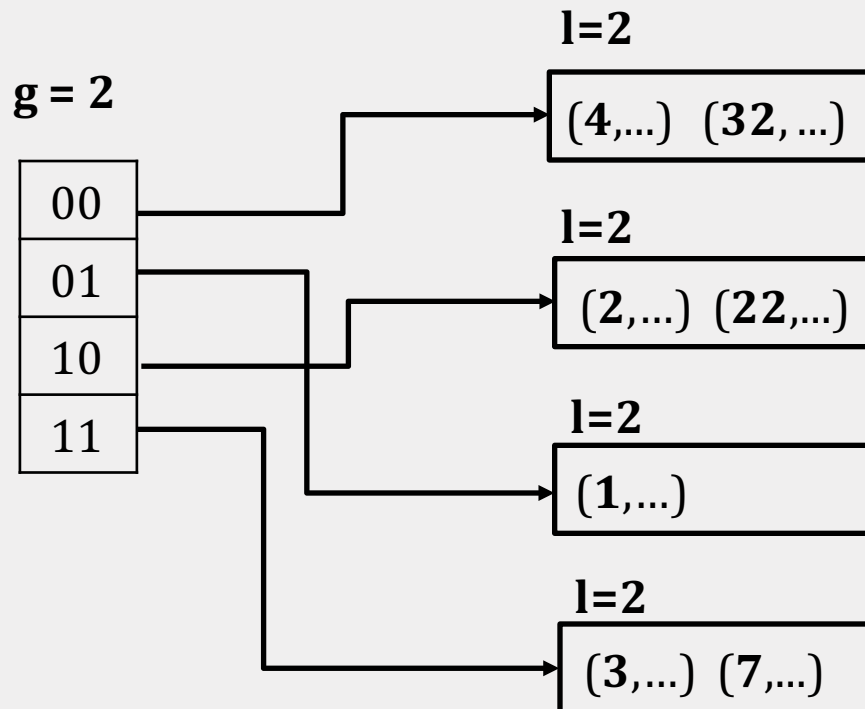
insert: (32,...)



# EXAMPLE

Since local depth is smaller than global,  
no need to change the directory size!

insert: (3,...)



The bucket is split into two

---

# EXTENDIBLE HASHING: DELETE

---

- Locate the bucket of the record and remove it
- If the bucket becomes empty, it can be removed (and update the directory)
- Two buckets can also be coalesced together if the sum of the entries fit in a single bucket
- Decreasing the size of the directory can also be done, but it is expensive

---

# MORE ON EXTENDIBLE HASHING

---

- How many disk accesses for equality search?
  - One if directory fits in memory, else two
- Directory grows in spurts, and, if the distribution of hash values is skewed, the directory can grow very large
- We may need overflow pages when multiple entries have the same hash value!