THE B+ TREE INDEX

CS 564- Fall 2018

WHAT IS THIS LECTURE ABOUT?

The **B+ tree** index

- Basics

- Search/Insertion/Deletion

- Design & Cost

INDEX RECAP

- We have the following query:
  ```
  SELECT *
  FROM Sales
  WHERE price > 100 ;
  ```

- How do we organize the file to answer this query efficiently?

INDEXES

- Hash index:
  - good for equality search
  - in expectation constant I/O cost for search and insert
- B+ tree index:
  - good for **range** and **equality** search

B+ TREE BASICS

THE B+ TREE INDEX

- a dynamic tree-structured index
  - adjusted to be always height-balanced
  - 1 node = 1 physical page
- supports efficient equality and range search
- widely used in many DBMSs
  - SQLite uses it as the default index
  - SQL Server, DB2, …

B+ TREE INDEX: BASIC STRUCTURE

root node

a node corresponds to a disk page

non-leaf nodes

leaf nodes

data entries
- exist *only* in the leaf nodes
- are sorted according to the search key

B+ TREE: NODE

- parameter d is the *order* of the tree
- each non-leaf node contains $d \leq m \leq 2d$ entries
  - minimum 50% occupancy at all times
- with the exception of the root node, which can have $1 \leq m \leq 2d$ entries

|	k_1		k_2		...		k_m	

NON-LEAF NODES

An non-leaf (or internal) node with m entries has $m+1$ pointers to lower-level nodes

	k_1		k_2		...		k_m	

pointer to a page with values $< k_1$

pointer to a page with $k_1 \leq$ values $< k_2$

pointer to a page with values $\geq k_m$

LEAF NODES

A leaf node with m entries has

- m pointers to the data records (rids)
- pointers to the next and previous leaves

	r_1	k_1	r_2	k_2	...	r_m	k_m	

pointer to the
previous page

pointer to the
next page

record   record   record

B+ Tree Operations

B+ TREE OPERATIONS

A B+ tree supports the following operations:

- equality search

- range search

- insert

- delete

- bulk loading

SEARCH

- start from the root node

- examine the index entries in non-leaf nodes to find the correct child

- traverse down the tree until a leaf node is reached
  - for equality search, we are done
  - for range search, traverse the leaves sequentially using the previous/next pointers
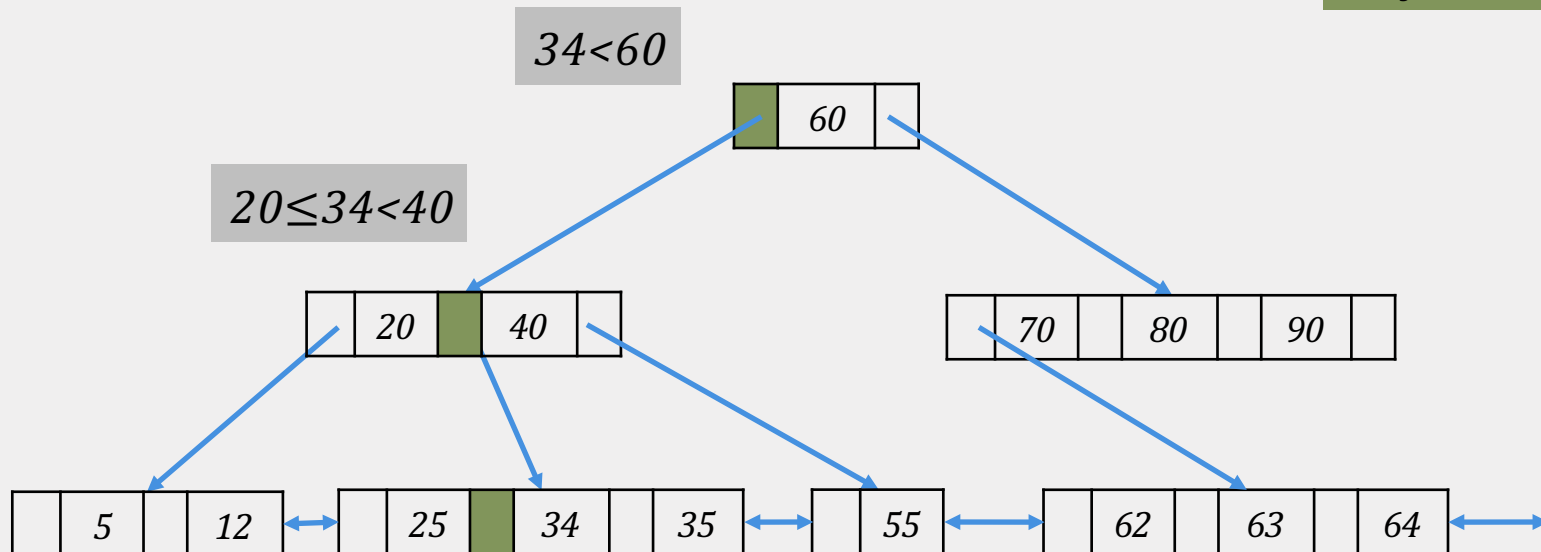
EQUALITY SEARCH: EXAMPLE

key = 34

34<60

| | 60 | |

| | 20 | | 40 | |

| | 70 | | 80 | | 90 | |

| | 5 | | 12 | ↔ | | 25 | | 34 | | 35 | ↔ | | 55 | ↔ | | 62 | | 63 | | 64 | ↔

EQUALITY SEARCH: EXAMPLE

key = 34

34<60

60

20≤34<40

| | 20 | | 40 | |

| | 70 | | 80 | | 90 | |

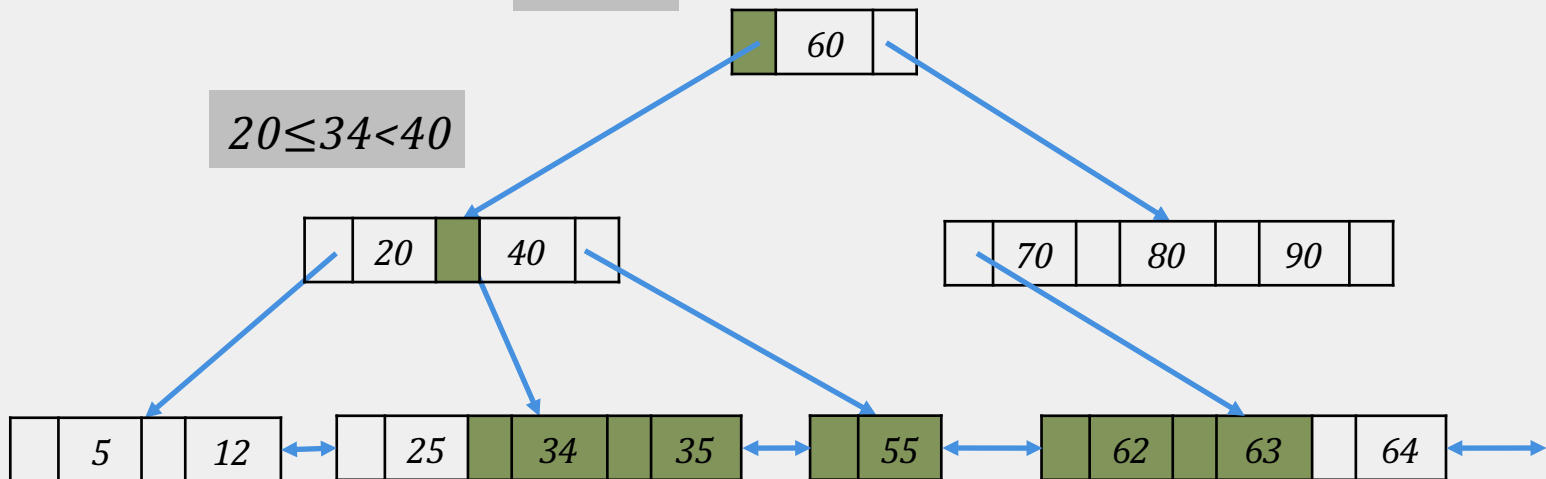| | 5 | | 12 | | ↔ | | 25 | | 34 | | 35 | ↔ | | 55 | ↔ | | 62 | | 63 | | 64 | ↔

To locate the correct data entry in the leaf node,
we can do either linear or binary search

RANGE SEARCH: EXAMPLE

34 ≤ key ≤ 63

34<60

20≤34<40

| | 60 | |

| | 20 | | 40 | |

| | 70 | | 80 | | 90 | |

| | 5 | | 12 | | → | | 25 | | 34 | 35 | → | | 55 | ← → | | 62 | 63 | | 64 | → |

After we find the leftmost point of the range,
we traverse sequentially!

INSERT

- find correct leaf node **L**
- insert data entry in **L**
  - If **L** has enough space, DONE!
  - Else, we must split **L** (into **L** and a new node **L'**)
    - redistribute entries evenly, **copy up** the middle key
    - insert index entry pointing to **L'** into parent of **L**
- This can propagate recursively to other nodes!
  - to split a non-leaf node, redistribute entries evenly, but **push up** the middle key

# INSERT: EXAMPLE

order **d** = 2

insert 8

| | 13 | | 17 | | 24 | | 30 | |

| 2 | 3 | 5 | 7 | 14 | 16 | 19 | 20 | 22 | 24 | 27 | 29 | 33 | 34 | 38 | 39 |

INSERT: EXAMPLE

order **d** = 2

insert 8

the leaf node is full so
we must split it!

| | 13 | | 17 | | 24 | | 30 | |

| 2 | 3 | 5 | 7 | | 14 | 16 | | 19 | 20 | 22 | | 24 | 27 | 29 | | 33 | 34 | 38 | 39 |

| 2 | 3 | | 5 | 7 | 8 |

d entries   d+1 entries

INSERT: EXAMPLE

order **d** = 2

the middle key (**5**) must be copied up, but the root node is full as well!

INSERT: EXAMPLE

order **d** = 2

insert 8

the new middle key (**17**) is now pushed up only (and not copied)

```
                          17

          5       13              24       30

  2  3    5  7  8    14  16    19  20  22    24  27  29    33  34  38  39
```

INSERT PROPERTIES

The B+ Tree insertion algorithm has several attractive qualities:

- ~ same cost as exact search

- it is ***self-balancing:*** the tree remains balanced (with respect to height) even after multiple insertions

B+ TREE: DELETE

- find leaf node **L** where entry belongs
- remove the entry
  - If **L** is at least half-full, DONE!
  - If **L** has only *d-1* entries,
    - Try to **re-distribute**, borrowing from sibling
    - If re-distribution fails, **merge L** and sibling
- If a merge occurred, we must delete an entry from the parent of **L**

DELETE : EXAMPLE 1

order **d** = 2

since by deleting 22 the node
remains half-full, we simply remove it

DELETE : EXAMPLE 1

order **d** = 2

delete 22

DELETE : EXAMPLE 2

order **d** = 2
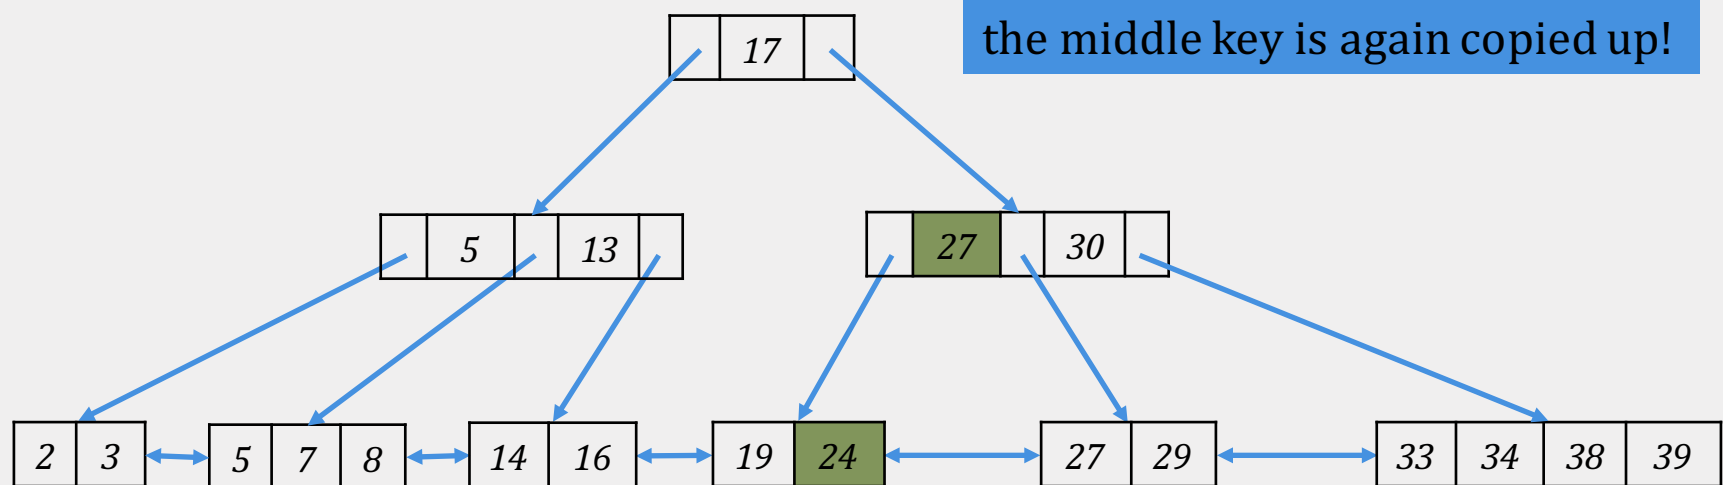
delete 20

by removing 20 the node is not half-full anymore, so we attempt to redistribute!

DELETE : EXAMPLE 2

order **d** = 2

delete 20

the middle key is again copied up!

```
                    ┌───┬────┬───┐
                    │ • │ 17 │ • │
                    └───┴────┴───┘
             ┌────────┘          └────────┐
      ┌───┬───┬───┬────┬───┐      ┌───┬────┬───┬────┬───┐
      │ • │ 5 │ • │ 13 │ • │      │ • │ 27 │ • │ 30 │ • │
      └───┴───┴───┴────┴───┘      └───┴────┴───┴────┴───┘
```

┌───┬───┐   ┌───┬───┬───┐   ┌────┬────┐   ┌────┬────┐   ┌────┬────┐   ┌────┬────┬────┬────┐
│ 2 │ 3 │ ↔ │ 5 │ 7 │ 8 │ ↔ │ 14 │ 16 │ ↔ │ 19 │ 24 │ ↔ │ 27 │ 29 │ ↔ │ 33 │ 34 │ 38 │ 39 │
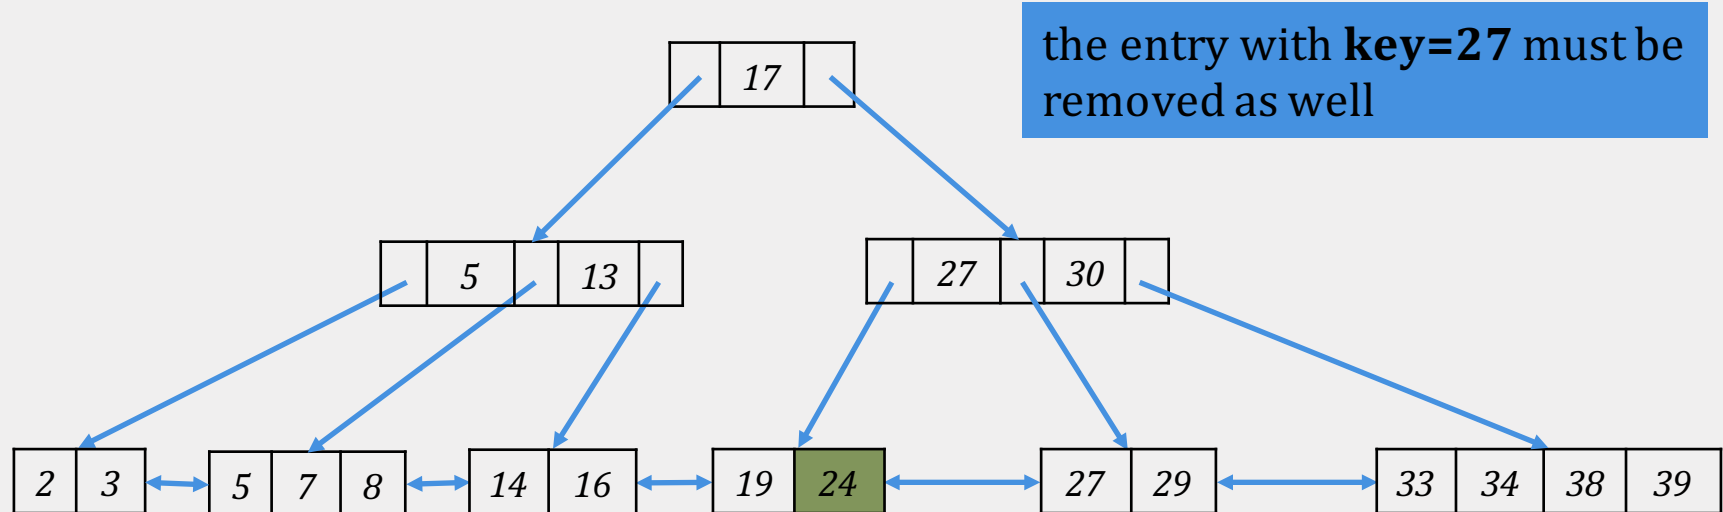└───┴───┘   └───┴───┴───┘   └────┴────┘   └────┴────┘   └────┴────┘   └────┴────┴────┴────┘

by removing 20 the node is not half-full anymore, so we attempt to redistribute!

DELETE : EXAMPLE 3

order **d** = 2

delete 24

the entry with **key=27** must be removed as well

17

5 | 13

27 | 30

2 | 3 ←→ 5 | 7 | 8 ←→ 14 | 16 ←→ 19 | **24** ←→ 27 | 29 ←→ 33 | 34 | 38 | 39

in this case, we have to merge nodes!

DELETE : EXAMPLE 3

order **d** = 2

delete 24

we are not done, since the resulting non-leaf node is not half-full!

DELETE : EXAMPLE 3

order **d** = 2

delete 24

we are not done, since the resulting non-leaf node is not half-full!

| - | 5 | - | 13 | - | 17 | | 30 | - |

| 2 | 3 | | 5 | 7 | 8 | | 14 | 16 | | 19 | 27 | 29 | | 33 | 34 | 38 | 39 |

B+ TREE: DELETE

- Redistribution of entries can also be possible for the non-leaf nodes

- We can also try to redistribute using *all siblings,* and not only the neighboring one

DUPLICATES

- **<u>duplicate keys</u>**: many data entries with the same key value
- Solution 1:
  - All entries with a given key value reside on a single page
  - Use overflow pages
- Solution 2:
  - Allow duplicate key values in data entries
  - Modify search operation

B+ TREE DESIGN & COST

B+ TREE: FAN-OUT

fan-out *f*: the number of pointers to child nodes coming out of a non-leaf node

- compared to binary trees (fan-out =*2*), B+ trees have a high fan-out ($d+1 \leq f \leq 2d+1$)
- The fan-out of B+ trees is dynamic, but we will often assume it is constant for our cost model

B+ TREE: FILL-FACTOR

<u>fill-factor</u> *F*: the percent of available slots in the B+ Tree that are filled

- it is usually < 1 to leave slack for (quicker) insertions!
- typical fill factor F = 2/3

B+ TREE: HEIGHT

height *h*: the number of levels of the non-leaf nodes

- the height is at least 1 (root node)
- high fan-out -> smaller height-> less I/O per search
- typical heights of B+ trees: 3 or 4

B+ TREE: EXAMPLE

- page size **P** = 4000 bytes

- search key size = 30 bytes

- address size = 10 bytes

- fill-factor **F** = 2/3

- number of records = 2,000,000

- We assume that the data entries store only the search key and the address of tuple

- We assume no duplicate entries

B+ TREE: EXAMPLE

What is the order **d** and fan-out **f** ?

- each non-leaf node stores up to 2**d** values of the key + (2**d**+1) addresses for the children pages
- to fit this into a single page, we must have:
$$2d \cdot 30 \ + (2d + 1) \cdot 10 \ \leq \ 4000$$
$$\boldsymbol{d} \leq 50$$
- since a maximum capacity node has (2**d**+1) = 101 children, and the fill-factor is 2/3, the fan-out is $\boldsymbol{f} = 101 \ * \frac{2}{3} = 67$

B+ TREE: EXAMPLE

How many leaf pages are in the B+ tree?

- we assume for simplicity that each leaf page stores only pairs of (key, address)
- each pair needs 30+10 = 40 bytes
- to store 2,000,000 such pairs with fill-factor **F** = 2/3, we need:

$$\#leaves = (2{,}000{,}000 * 40)/(4{,}000 * \textbf{F}) = 30{,}000$$

B+ TREE: EXAMPLE

What is the height **h** of the B+ tree?

- we calculated that we need to index **N** = 30,000 pages
- **h** = 1 -> indexes **f** pages
- **h** = 2 -> indexes **f**^{2} pages
- ...
- **h** = k -> indexes **f**^{k} pages

height must be $\boldsymbol{h} = \lceil log_f N \rceil$

for our example, $\boldsymbol{h} = \lceil log_{67} 30{,}000 \rceil = 3$

B+ TREE: EXAMPLE

What is the total size of the tree?

- #pages = 1 + 67 + 67^2 + 30,000 = 34,557
- the top levels of the B+ tree do not take much space and can be kept in the buffer pool
  - *level 0 =*   1 page   ~   4 KB
  - *level 1 =*   67 pages ~   268 KB
  - *level 2 =*   4,489 pages ~   18 MB

COST MODEL FOR SEARCH

To do equality search:

- we read one page per level of the tree
- levels that we can fit in buffer are free!
- finally we read in the actual record

$$\text{I/O cost } = \boldsymbol{h} - L_B + 1 + I$$

$I = 0$ if the record is stored at the leaf node, otherwise $I = 1$

If we have $\boldsymbol{B}$ available buffer pages, we can store $\boldsymbol{L}_B$ levels of the B+ Tree in memory:

- $\boldsymbol{L}_B$ is the number of levels such that the sum of all the levels' nodes fit in the buffer:

$$B \geq 1 + f + \cdots + f^{L_B - 1}$$

COST MODEL FOR SEARCH

To do range search:

- we read one page per level of the tree
- levels that we can fit in buffer are free!
- we read sequentially the pages in the range

$$\text{I/O cost } = \text{h} - L_B + OUT$$

Here, OUT is the I/O cost of loading the additional leaf nodes we need to access + the I/O cost of loading each **page** of the results