

# BIG DATA SYSTEMS

---

*CS 564- Fall 2018*

---

*ACKs: Magda Balazinska*

# WHAT IS BIG DATA?

---

The three V's:

- high **Volume**
- high **Variety**
- high **Velocity**

# VOLUME

---

- Databases parallelize easily
- Techniques available from the 80's (GAMMA project started @ Wisconsin!)
  - data partitioning
  - parallel query processing
- SQL is **embarrassingly parallel**

# VARIETY

---

- complex workloads:
  - Machine Learning tasks: e.g. click prediction, topic modeling, SVM, k-means
- various types of data:
  - text data
  - semi-structured data
  - graph data
  - multimedia (video, photos)

# VELOCITY

---

- data is generated very fast
- data needs to be processed very fast
- real time data analytics
- **data streaming** (each data item can be processed only once!): e.g. financial data

---

# ANOTHER V: VERACITY

---

The data collected is often **uncertain**

- inconsistent data (violated constraints)
- incomplete data (missing values)
- ambiguous data

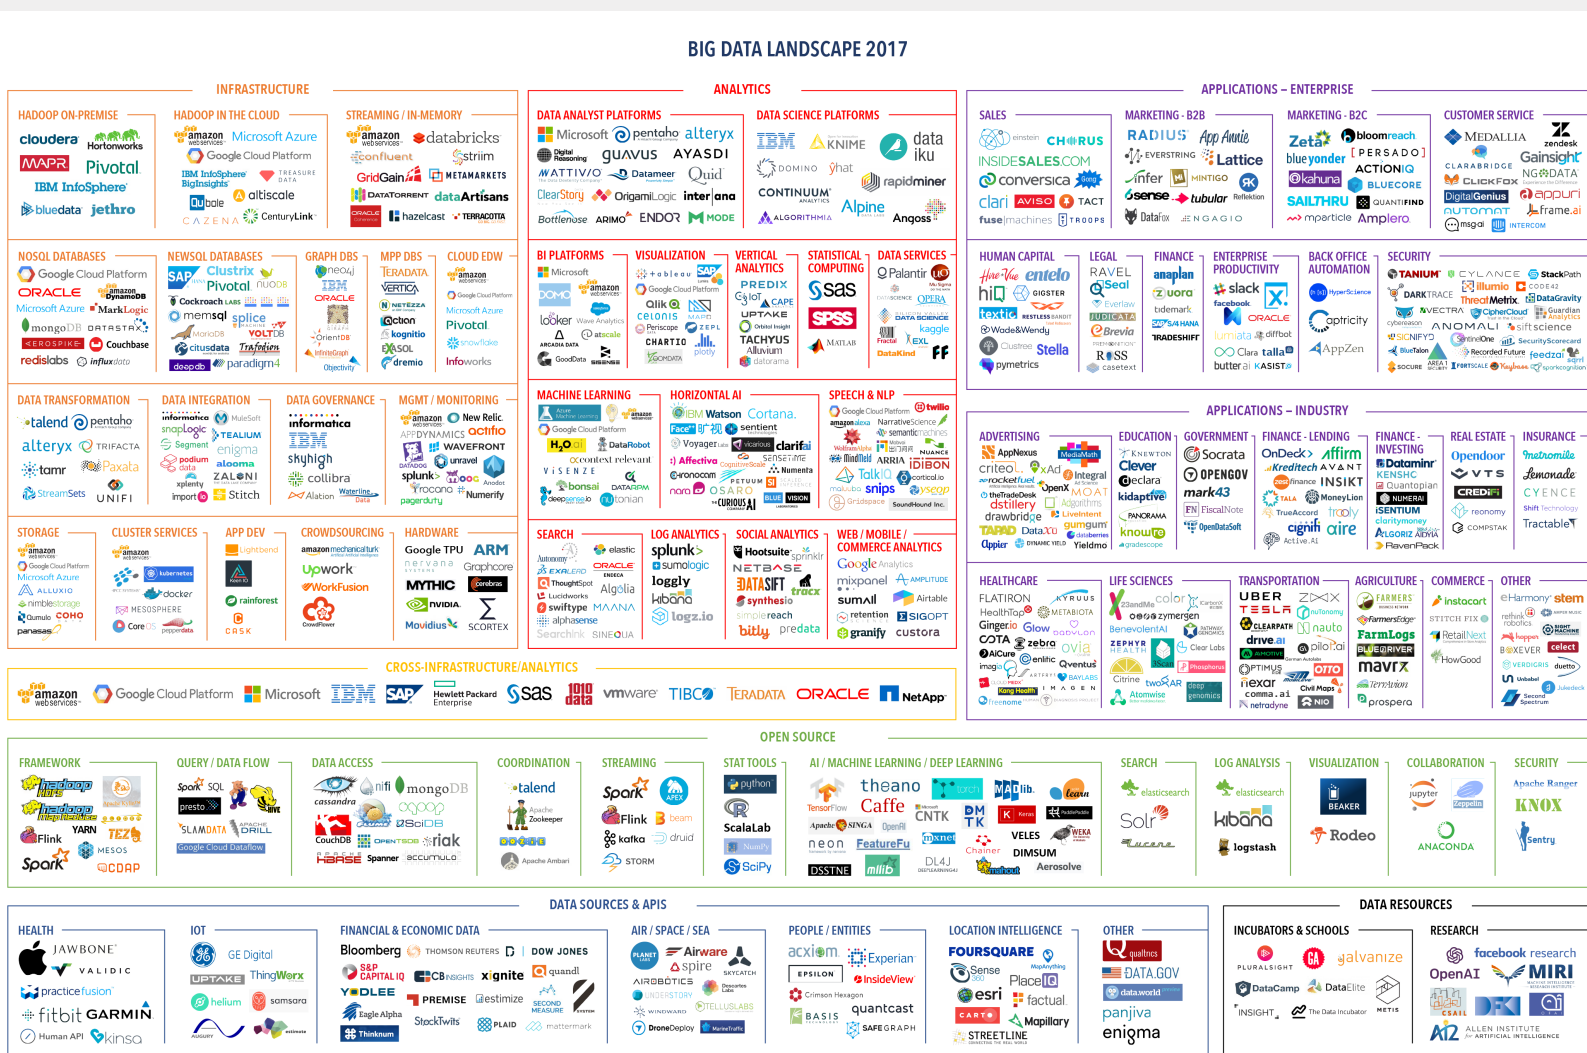
Example: sensor data

---

# DATA ANALYTICS COMPANIES

---

- **Greenplum**: founded in 2003 acquired by EMC in 2010. A parallel shared-nothing DBMS
- **Vertica**: founded in 2005 and acquired by HP in 2011. A parallel column-store shared-nothing DBMS
- **AsterData**: founded in 2005 acquired by Teradata in 2011. A parallel, shared-nothing, MapReduce-based data processing system
- **Netezza**: founded in 2000 and acquired by IBM in 2010. A parallel shared-nothing DBMS





---

# SCALING FOR BIG DATA

---

- **OLTP:** Online Transaction Processing
  - **scale** transactions per second
  - Amazon, Facebook, Twitter, ...
- **OLAP:** Online Analytical Processing
  - **scale** query response time
  - analysis of massive datasets, decision making, ...

## 2 APPROACHES

---

- Parallel databases, started at the 80s. Both for
  - OLTP (transaction processing)
  - OLAP (decision support queries)
- MapReduce
  - first developed by Google, published in 2004
  - only for decision support queries
  - ecosystem around it: Hadoop, PigLatin, Hive, ...

Today these two approaches have been converging!

---

# PARALLEL DBMS

---

- The goal is to improve performance by executing multiple operations in parallel
- Terminology to measure performance:
  - **Speed-up**: using more processors, how much faster does the task run (if problem size is fixed)?
  - **Scale-up**: using more processors, does performance remain the same as we increase the problem size?

# SCALE-UP VS SCALE-OUT

---

## Scale-up

- using more powerful machines, more processors/RAM per machine

## Scale-out

- using a larger number of servers

# ARCHITECTURES

---

- Shared memory
  - nodes share RAM + disk
  - easy to program, expensive to scale
- Shared disk
  - nodes access the same disk, hard to scale
- Shared nothing
  - nodes have their own RAM+disk
  - connected through a fast network

---

# PARALLEL QUERY EVALUATION

---

- Multiple DBMS instances also called **nodes** execute on machines in a cluster
  - one instance plays role of the **coordinator**
  - other instances play role of **workers**
- Workers execute queries
  - typically all workers execute the same plan (intra-operator parallelism, intra-query parallelism)
  - workers can execute multiple queries at the same time (inter-query parallelism)

---

# PARALLEL DATA STORAGE

---

Horizontal data partitioning

- **block** partitioned
- **hash** partitioned
- **range** partitioned

Uniform vs skewed partitioning

---

# PARALLEL QUERY ALGORITHMS

---

- Parallel Selection
- Parallel Join
  - hash join
  - broadcast join



---

# BEYOND PARALLEL DBMS

---

Many big data systems go beyond the capabilities of relational systems:

- MapReduce/Hadoop
- Spark, SparkSQL
- Mlib, Tensorflow
- Graph processing systems
- Stream processing systems