


Fall 2019

LOGGING AND RECOVERY

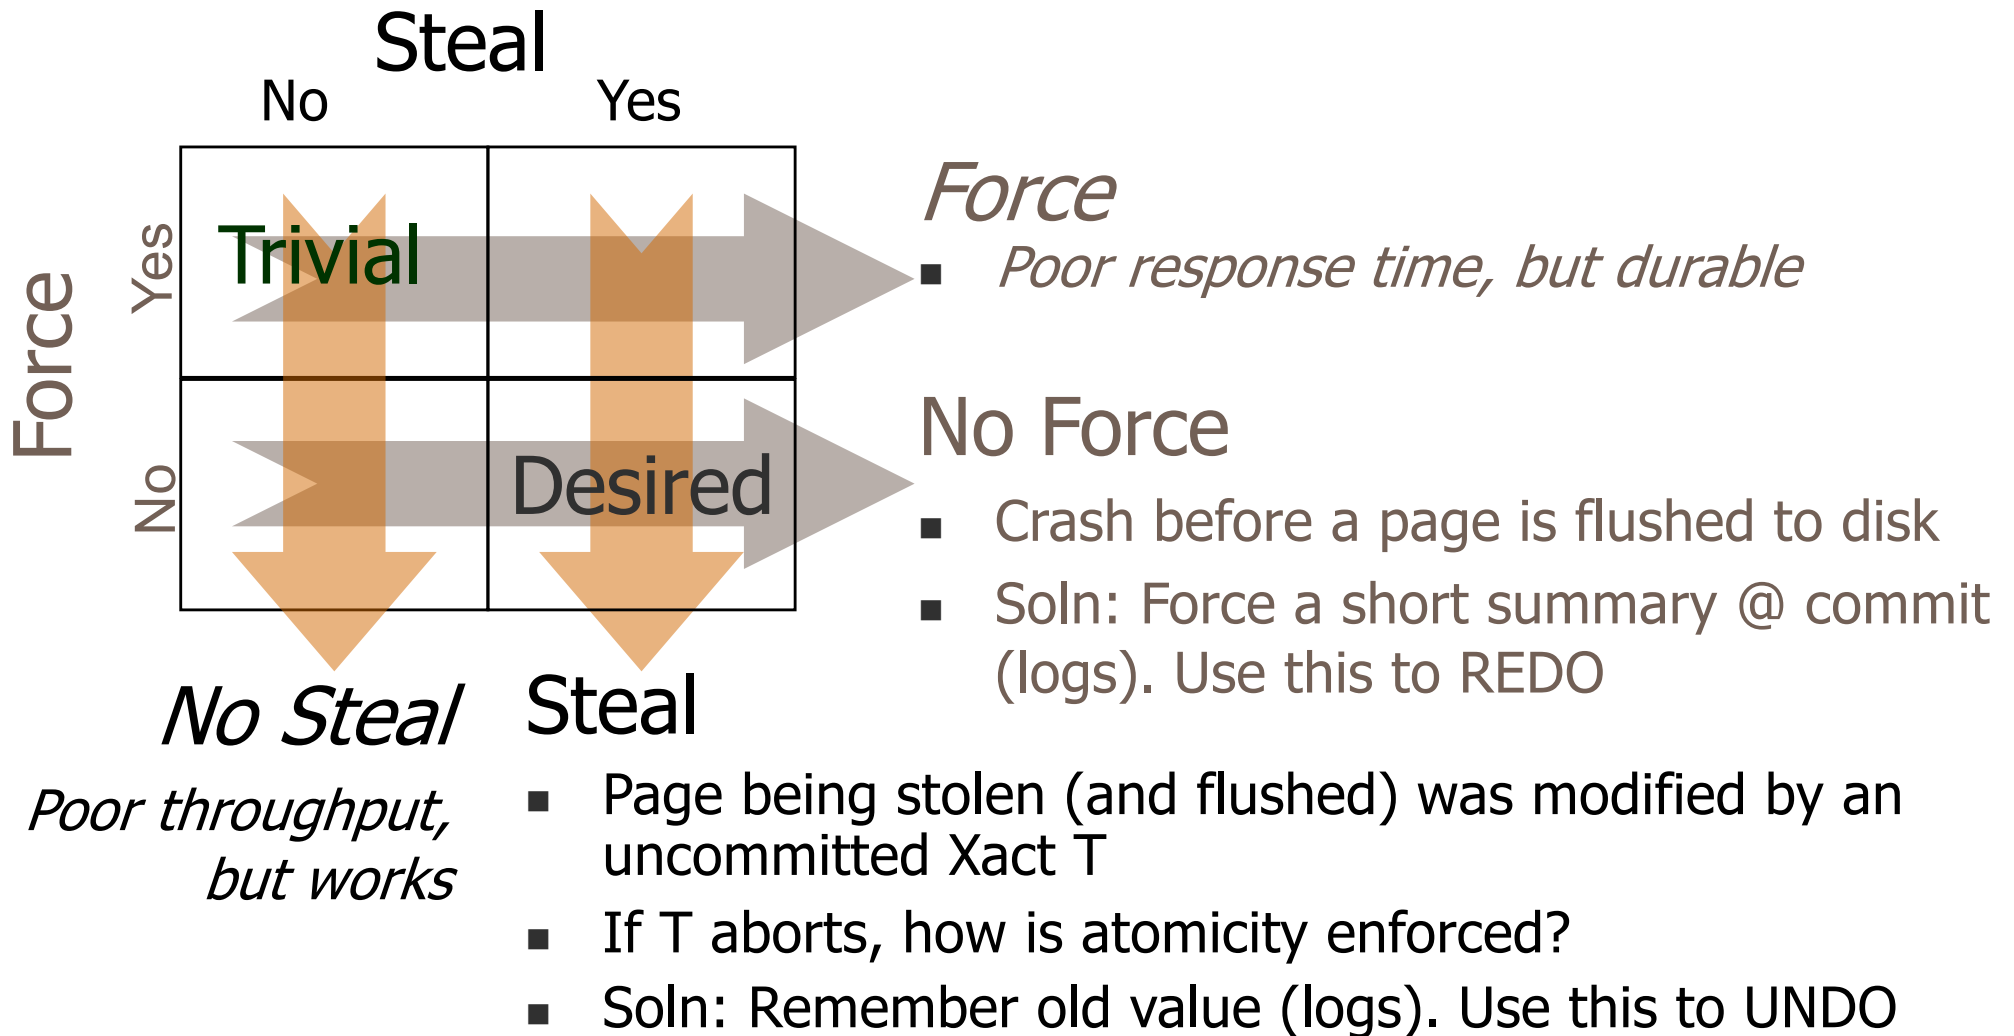
[CH 18]

Motivation

- Atomicity:
 - Transactions may abort (“Rollback”).
- Durability:
 - What if DBMS stops running? (Causes?)
- Desired Behavior after system restarts:
 - T1, T2 & T3 should be durable.
 - T4 & T5 should be aborted



Buffer Pool: Sharing & Writing




Basic Idea: Logging

- Record information, for every change, in a *log*.
 - Sequential writes to log (put it on a separate disk).
 - Stored in stable storage to survive system crash
 - disk mirroring
 - Each record has a log sequence number (LSN)
 - Log record contains:
 - `<prevLSN, XID, type, ... >`
 - and additional control info (which we'll see soon)
 - Note: the log records for a transaction are chained by prevLSN

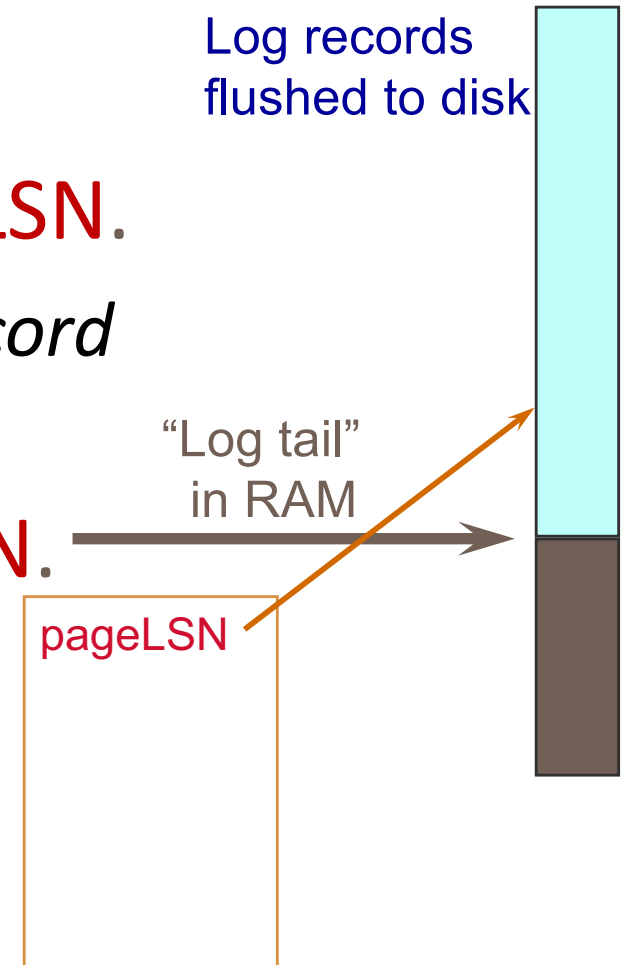
Write-Ahead Logging (WAL)

- The **Write-Ahead Logging** Protocol:
 1. Must force the log record for an update before the corresponding data page gets to disk.
 2. Must write all log records for a Xact before commit.
- #1 guarantees Atomicity.
- #2 guarantees Durability.
- Exactly how is logging (and recovery!) done?
 - We'll study the ARIES algorithms
 - breakthrough in recovery algorithms!
 - repeating history paradigm
 - fine-granularity locking and logical logging

WAL & the Log



- Log Sequence Number (LSN).
 - Unique and always increasing.
- Each *data page* contains a **pageLSN**.
 - The LSN of the most recent *log record* that updated the page
- System keeps track of **flushedLSN**.
 - The max LSN flushed so far.
- WAL: *Before* a page is written,
 - $\text{pageLSN} \leq \text{flushedLSN}$



Log Records


LogRecord fields:

prevLSN
XID
type
pageID
length
offset
before-image
after-image
...

update
records
only

Possible log record types:

- **Update**
- **Commit**
- **Abort**
- **End** (end of commit or abort)
- **Compensation Log Rec. (CLRs)**
 - For UNDO actions. When?
 - Contains **undoNextLSN**
 - Reverse chain of update logs
 - Contains before-image



Other Log-Related State

- **Transaction Table**: One entry per active Xact.
 - Contains
 - **XID**: Transaction identifier
 - **status**: running/committed/aborted
 - **lastLSN**: LSN of the most recent log rec. for this Xact.
- **Dirty Page Table**: One entry per dirty page in BP
 - Contains **recLSN**:- LSN of the log record that ***first*** caused the page to be dirty.
 - Starting point for REDO

Checkpointing

- **Checkpoint**: Snapshot of the database
 - Minimize recovery time
- Write to log:
 - **begin_checkpoint** record: Indicates when chkpt began.
 - **end_checkpoint** record:
 - Record Xact table and D.P.T.
 - Tables accurate only as of the time of the **begin_checkpoint** record
 - No attempt to force dirty pages to disk
 - This is a **fuzzy checkpoint**
 - Store LSN of chkpt record in a safe place (**master** record).

Normal Execution of an Xact


- Series of **reads** & **writes**, followed by **commit** or **abort**.
 - Updates are “in place”: i.e., data on disk is overwritten
 - We will assume that write is atomic on disk.
 - In practice, additional details to deal with non-atomic writes.
- **Strict 2PL**.
- STEAL, NO-FORCE buffer management, with **Write-Ahead Logging**.

The Big Picture: What's Stored Where




LogRecords

prevLSN
XID
type
pageID
length
offset
before-image
after-image



Data pages
each
with a
pageLSN

master record



Xact Table

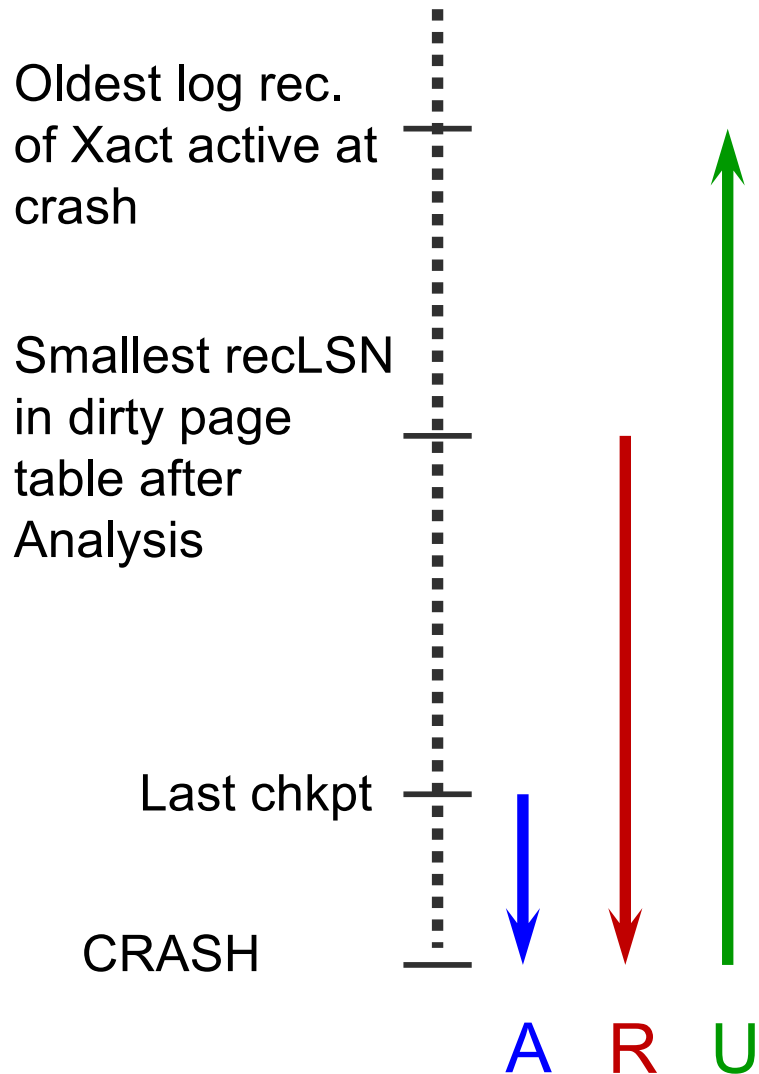
lastLSN
status

Dirty Page Table

recLSN

flushedLSN

Crash Recovery: Big Picture

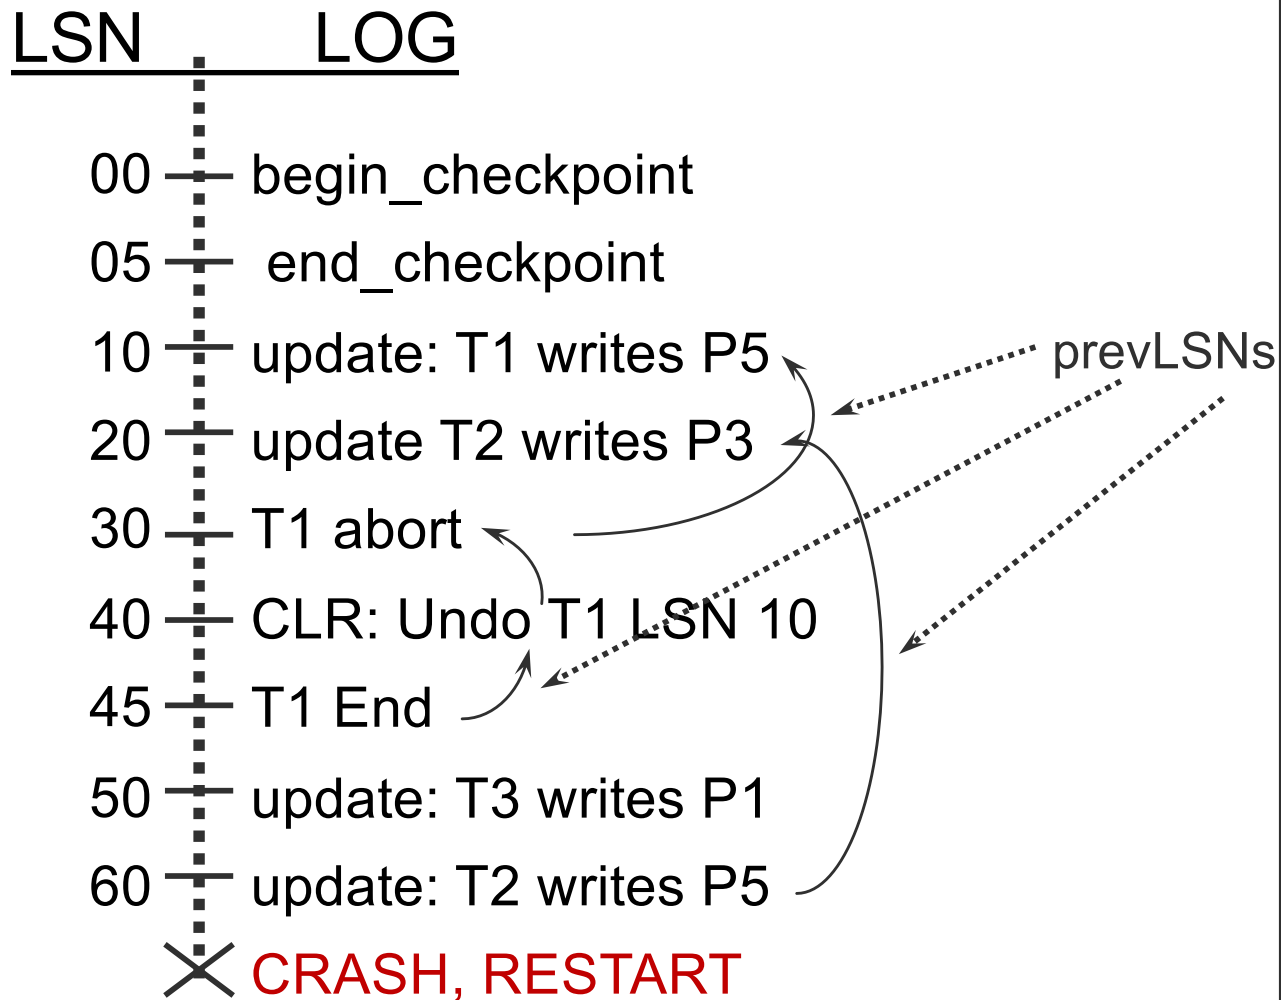


- Start from a **checkpoint** (found via **master** record).
- Three phases:
 - **Analysis**: Since checkpoint, find
 - Xacts that must be aborted (losers)
 - dirty pages at time of crash (conservative estimate)
 - **REDO** *all* actions.
 - repeat history
 - **UNDO** effects of losers

Recovery: The Analysis Phase

- Compute
 - Set of dirty pages (conservative)
 - Uncommitted transactions at the crash point
- Scan log forward from checkpoint.
 - End record: Remove Xact from Xact table.
 - Other records: Add Xact to Xact table, set lastLSN=LSN
 - Commit record: change Xact status to commit.
 - Update or CLR record: If P not in Dirty Page Table,
 - Add P to D.P.T., set its recLSN=LSN.

Recovery: The Analysis Phase



Transaction Table

XACT	LastLSN
T2	60
T3	50

D. P. T.

Page	recLSN
P1	50
P3	20
P5	10

Recovery: The REDO Phase

- *Repeat History* to reconstruct state at crash:
 - Reapply *all* updates (even of aborted Xacts!), redo CLR
 - Bring the database to the same state as @ crash
- Scan forward from log record containing smallest *recLSN* in D.P.T.

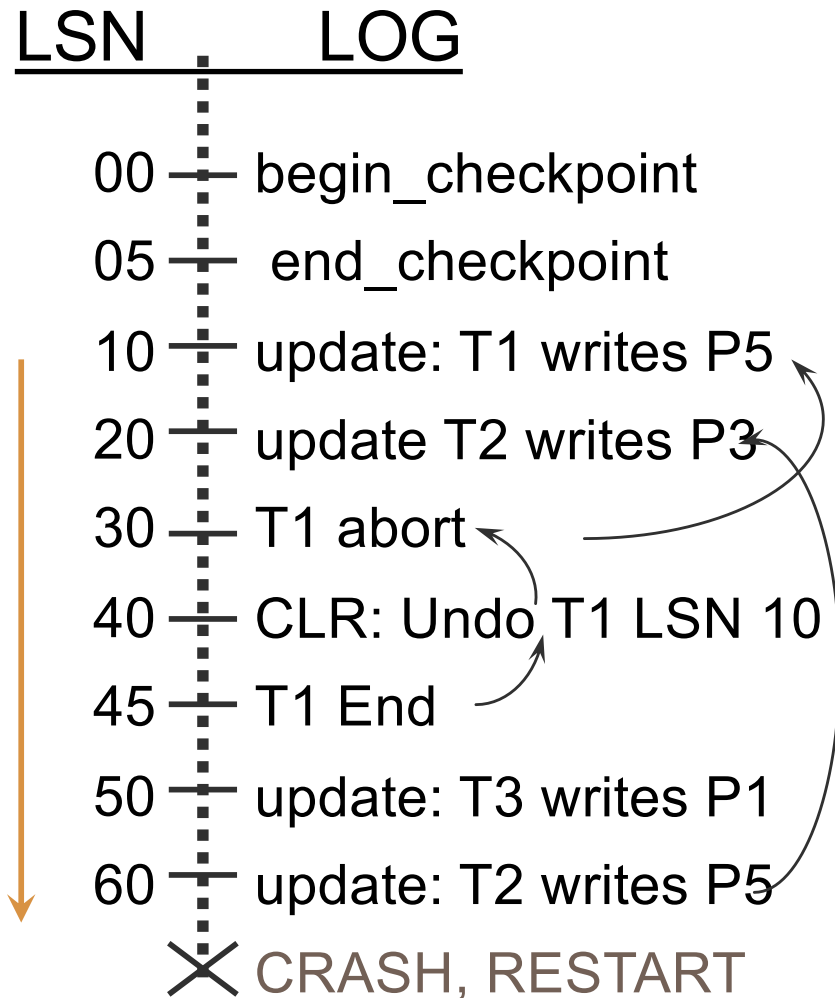
For each update log record or CLR, REDO the action unless we can verify that the change has already been written to disk:

- Affected page is not in the Dirty Page Table, or
- Affected page is in D.P.T., but $LSN < recLSN$, or
 - update was propagated to disk
- $LSN \leq pageLSN$ (in DB)
 - requires fetching the page

To REDO An Action

- Reapply logged action.
- Set pageLSN to LSN. No additional logging!
- Use of CLRs ensures that no change is ever carried out twice on the disk copy of an object.
 - For every “DO” there is one and only one “UNDO”
- At the end of REDO
 - Write END log recs for all committed Xacts.
 - Remove committed Xacts from the Xact table.

Recovery: The REDO Phase



Transaction Table

XACT	LastLSN
T2	60
T3	50

D. P. T.

Page	recLSN
P1	50
P3	20
P5	10

Recovery: The UNDO Phase

ToUndo = {lastLSNs of all “loser” Xact}

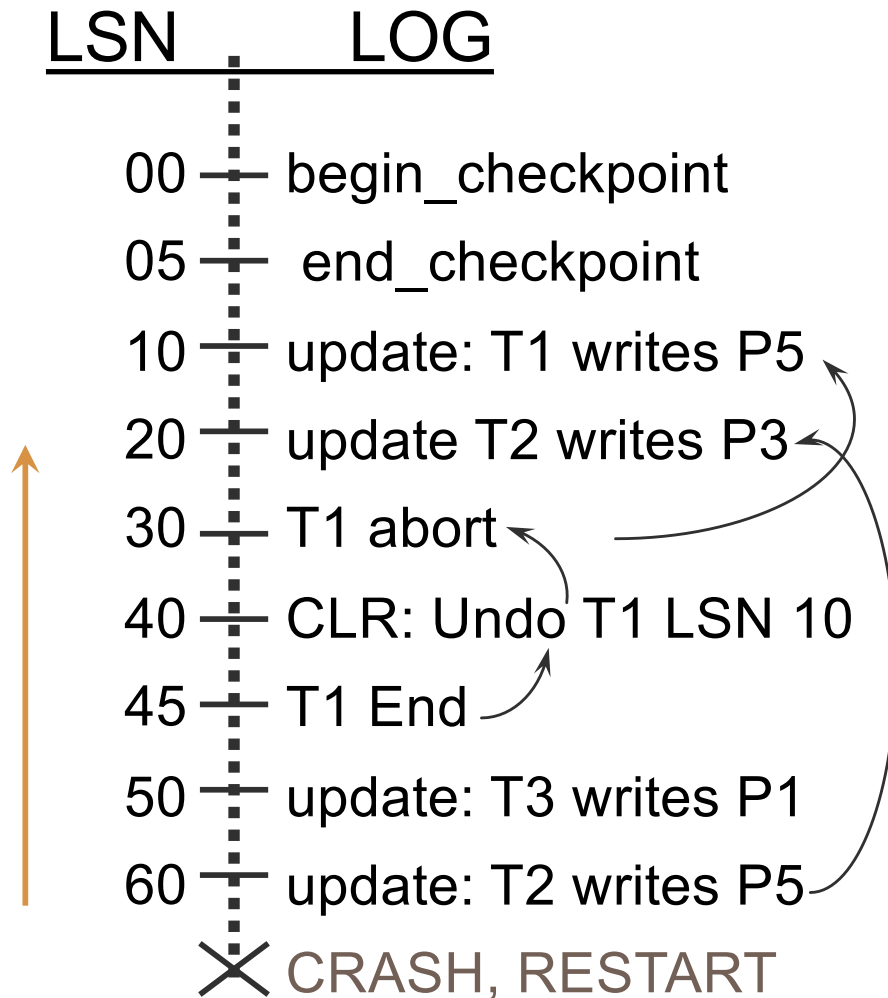
Repeat:

- Choose largest LSN among ToUndo.
- If this LSN is a CLR and undonextLSN==NULL
 - Write an End record for this Xact.
- If this LSN is a CLR, and undonextLSN != NULL
 - Add undonextLSN to ToUndo
- Else this LSN is an update. Undo the update, write a CLR, add prevLSN to ToUndo.

Abort: special case of UNDO

Until ToUndo is empty.

Recovery: The UNDO Phase



XACT	LastLSN
T2	60
T3	50

Page	recLSN
P1	50
P3	20
P5	10

LSN	LOG	(undoNextLSN)
70	CLR: Undo T2, LSN 60,	(20)
80	CLR: Undo T3, LSN 50,	(null)
85	T3 End	
90	CLR: Undo T2, LSN 20,	(null)
95	T2 End	

Example: Crash During Restart!

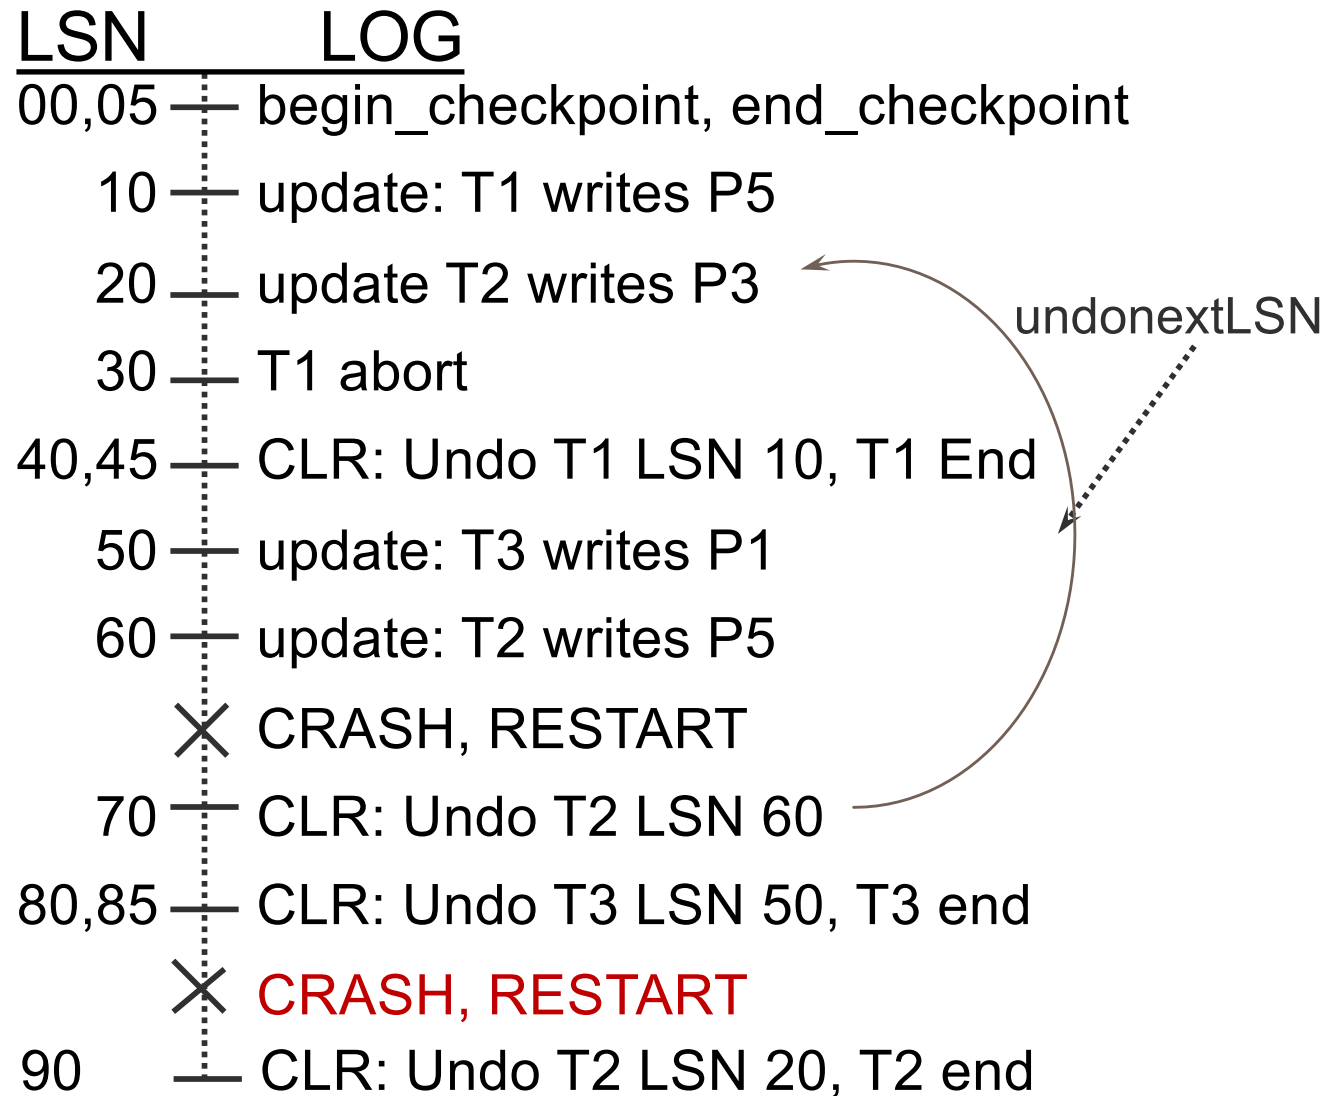
XACT	LastLSN
T2	70

Page	recLSN
P1	50
P3	20
P5	10

REDO: 10 to 85

UNDO:

- Undo 70, CLR
- Undo 20
- Take a ckpt



Additional Crash Issues

- How do you limit the amount of work in REDO?
 - Flush asynchronously in the background.
 - Watch “hot spots”!
- How do you limit the amount of work in UNDO?
 - Avoid long-running Xacts.

Media Recovery

- Used for disaster recovery.
- Based on periodically making a copy of the database
 - similar to a fuzzy checkpoint
- Apply logs to the copy of the object in the media to bring it up-to-date

Summary of Logging/Recovery

- Atomicity & Durability.
- WAL to allow STEAL/NO-FORCE
- Checkpointing: A quick way to limit the amount of log to scan on recovery.
- Recovery works in 3 phases:
 - Analysis: Forward from checkpoint.
 - Redo: Forward from oldest recLSN.
 - Undo: Backward from end to first LSN of oldest Xact alive at crash.
- Upon Undo, write CLR.
- Redo “repeats history”: Simplifies the logic!
- **Interested in the history of ARIES:**
 - <http://www.sigmod.org/publications/ds-collection/discs-new/2000.1/out/slides/vldb/repeatinghistoryc/index.pdf>