


1c.



3f.

```

29 void mmul3(const double* A, const double* B_T, double* C, const std::size_t n)
30 {
31     int len = (int)n;
32     for (int A_row_iter = 0; A_row_iter < len; A_row_iter++) {
33         for (int B_col_iter = 0; B_col_iter < len; B_col_iter++) {
34             C[A_row_iter*n + B_col_iter] = 0;
35             for (int k = 0; k < len; k++) {
36                 C[A_row_iter*n + B_col_iter] += A[A_row_iter*n + k] * B_T[B_col_iter*n + k];
37             }
38         }
39     }
40 }
41
42 void mmul4(const double* A_T, const double* B, double* C, const std::size_t n)
43 {
44     int len = (int)n;
45     for (int A_row_iter = 0; A_row_iter < len; A_row_iter++) {
46         for (int B_col_iter = 0; B_col_iter < len; B_col_iter++) {
47             C[A_row_iter*n + B_col_iter] = 0;
48             for (int k = 0; k < len; k++) {
49                 C[A_row_iter*n + B_col_iter] += A_T[A_row_iter + k*n] * B[B_col_iter + k*n];
50             }
51         }
52     }
53 }
54

```

Since there is cache in hardware, we could utilize cache to speed up program execution time by either utilizing spatial locality or temporal locality. mmul3 is such a function that uses spatial locality in its most inner loop (line 36) which do a stride 1 increment when k increases by 1. On the contrary, mmul4 do a stride n increment when k increases by 1 (line 49), so mmul4 is less cache-friendly. As a result, mmul3 is much quicker than mmul4.